

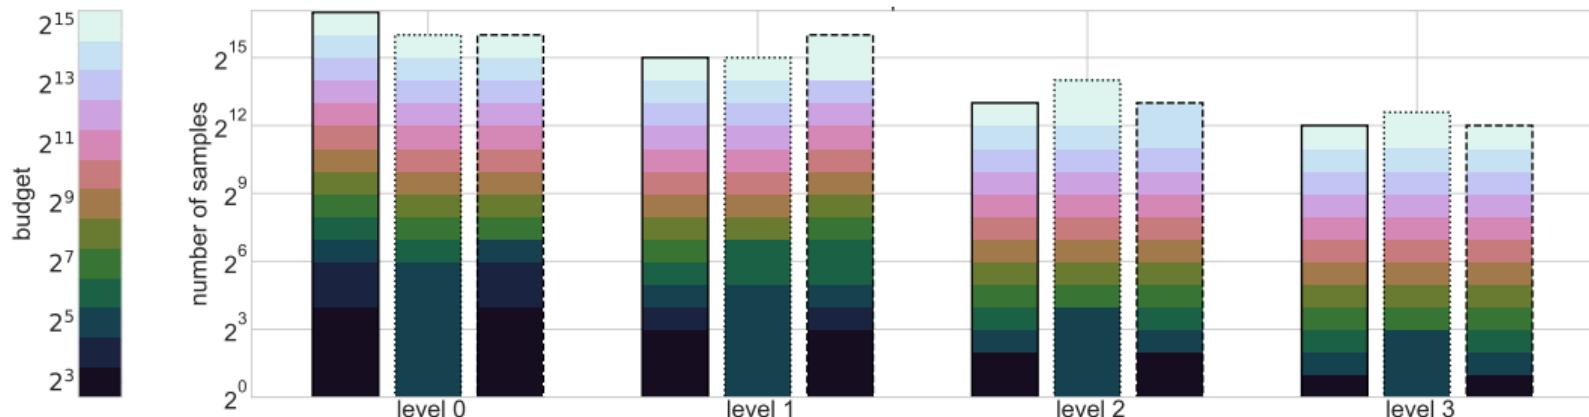
Fast Bayesian Multilevel Quasi-Monte Carlo

Aleksei G. Sorokin^{1,2}, Pieterjan Robbe², Gianluca Geraci²
Michael S. Eldred², Fred J. Hickernell¹,

¹Illinois Institute of Technology, Department of Applied Mathematics.

² Sandia National Lab.

[1] <https://arxiv.org/abs/2510.24604>



Motivating Example

2D Nonlinear Darcy flow PDE with random forcing term g and conductivity e^a

$$\begin{cases} -\nabla \cdot (e^{a(\mathbf{u}, \mathbf{X})} \cdot \nabla q(\mathbf{u}, \mathbf{X})) + (q(\mathbf{u}, \mathbf{X}))^3 = g(\mathbf{u}, \mathbf{X}), & \mathbf{u} \in [0, 1]^2 \\ q(\mathbf{u}, \mathbf{X}) = 0, & \mathbf{u} \in \partial[0, 1]^2. \end{cases}$$

Coefficients (g, a) are independent Gaussian processes with stochasticity $\mathbf{X} \sim \mathcal{U}[0, 1]^d$.
 $q_\ell(\mathbf{u}, \mathbf{X})$ a numerical PDE solution on a computational grid with $\mathcal{O}(2^{2\ell})$ nodes

- level $\ell \in \{1, \dots, L\}$ controls solution accuracy: larger $\ell \rightarrow$ finer mesh
- Numerical PDE solution from a Levenberg–Marquardt iteration [2, 3]
- Stochastic dimension $d = \mathcal{O}(2^{2L})$, nodes in the finest mesh (the largest level)

Quantity of interest

$$\mathbb{E}[Q_L(\mathbf{X})] := \mathbb{E} \left[\max_{\mathbf{u} \in [0, 1]^2} q_L(\mathbf{u}, \mathbf{X}) \right], \quad \mathbf{X} \sim \mathcal{U}[0, 1]^d$$

2D Nonlinear Darcy Flow PDE

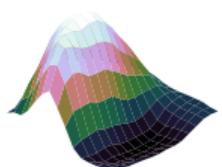
$$\mathbb{E}[Q_L(\mathbf{X})] := \mathbb{E} \left[\max_{\mathbf{u} \in [0,1]^2} q_L(\mathbf{u}, \mathbf{X}) \right], \quad \mathbf{X} \sim \mathcal{U}[0,1]^d$$

solution $q_0(\mathbf{u}, \mathbf{x})$
 $Q_0(\mathbf{x}) = 0.0288$



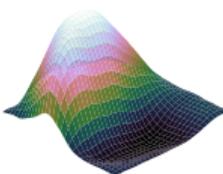
$Q_0(\mathbf{x}) = 0.0302$

solution $q_1(\mathbf{u}, \mathbf{x})$
 $Q_1(\mathbf{x}) = 0.0374$



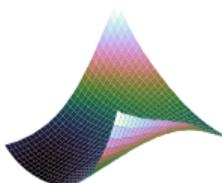
$Q_1(\mathbf{x}) = 0.0281$

solution $q_2(\mathbf{u}, \mathbf{x})$
 $Q_2(\mathbf{x}) = 0.0369$

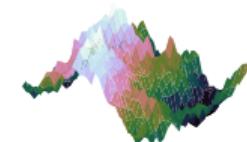
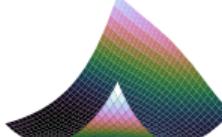
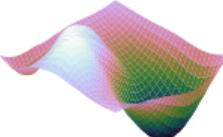
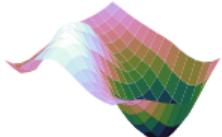


$Q_2(\mathbf{x}) = 0.0349$

conductivity $e^{a(\mathbf{u}, \mathbf{x})}$



forcing term $g(\mathbf{u}, \mathbf{x})$



The Multilevel Monte Carlo Problem

Can we accelerate high-fidelity UQ by exploiting cheaper low-fidelity surrogates?

$Q_\ell : [0, 1]^d \rightarrow \mathbb{R}$ simulates at (fidelity) level $\ell \in \{1, \dots, L\}$ with cost C_ℓ increasing in ℓ .

- e.g., Q_ℓ a (possibly nonlinear) functional of numerical PDE solution q_ℓ

Stochasticity identified with $\mathbf{X} \sim \mathcal{U}[0, 1]^d$ (after a possible transform)

- e.g., \mathbf{X} used to generate random PDE coefficients
- d may equal number of nodes in the computational grid
- d may equal number of coefficients in truncated Karhunen–Loève expansion
- d may change with level, e.g., d_ℓ the number of monitoring times in option pricing

Quantity of interest is the expectation (mean) at the highest-level (highest-fidelity)

$$\nu = \mathbb{E}[Q_L(\mathbf{X})] = \sum_{\ell=1}^L \mathbb{E}[Q_\ell(\mathbf{X}) - Q_{\ell-1}(\mathbf{X})] = \sum_{\ell=1}^L \mathbb{E}[Y_\ell(\mathbf{X})] = \sum_{\ell=1}^L \mu_\ell$$

where $Y_\ell := Q_\ell - Q_{\ell-1}$, $\mu_\ell := \mathbb{E}[Y_\ell]$, and we set $Q_0 = 0$ so $Y_1 = Q_1$.

Multilevel Monte Carlo (MLMC) with IID Points [4]

$$\nu = \mathbb{E}[Q_L(\mathbf{X})] = \sum_{\ell=1}^L \mathbb{E}[Q_\ell(\mathbf{X}) - Q_{\ell-1}(\mathbf{X})] = \sum_{\ell=1}^L \mathbb{E}[Y_\ell(\mathbf{X})] = \sum_{\ell=1}^L \mu_\ell, \quad \mathbf{X} \sim \mathcal{U}[0,1]^d.$$

Approximate ν by the sum of n_ℓ -sample Monte Carlo approximations

$$\hat{\nu} = \sum_{\ell=1}^L \hat{\mu}_\ell \quad \text{where} \quad \hat{\mu}_\ell = \frac{1}{n_\ell} \sum_{i=0}^{n_\ell-1} Y_\ell(\mathbf{x}_{\ell,i}),$$

using IID points $\mathbf{x}_{1,0}, \dots, \mathbf{x}_{1,n_1-1}, \dots, \mathbf{x}_{L,0}, \dots, \mathbf{x}_{L,n_L-1} \stackrel{\text{IID}}{\sim} \mathcal{U}[0,1]^d$.

$\therefore \mathbb{V}[\hat{\nu}] = \sum_{\ell=1}^L V_\ell/n_\ell$ for variance $V_\ell = \mathbb{V}[Y_\ell(\mathbf{X})]$.

\therefore for a fixed error tolerance, optimal allocation $n_\ell \propto \sqrt{V_\ell/C_\ell}$.

Algorithm 1 MC: MLMC with IID Points

Require: $N, C_\ell, n_\ell^{\text{next}}$ \triangleright budget, costs, initial samples

$n_\ell \leftarrow 0$ for $\ell \in \{1, \dots, L\}$ and $\mathcal{L} \leftarrow \{1, \dots, L\}$

while true **do**

 Generate $\mathbf{x}_{\ell,i} \sim \mathcal{U}[0,1]^d$ for $\ell \in \mathcal{L}$ and $n_\ell \leq i < n_\ell^{\text{next}}$ all independently

 Evaluate $Y_\ell(\mathbf{x}_{\ell,i})$ for $\ell \in \mathcal{L}$ and $n_\ell \leq i < n_\ell^{\text{next}}$

$\hat{\mu}_\ell \leftarrow \frac{1}{n_\ell} \sum_{i=0}^{n_\ell-1} Y_\ell(\mathbf{x}_{\ell,i})$ for $\ell \in \mathcal{L}$ \triangleright estimate of μ_ℓ

$\sigma_\ell^2 \leftarrow \frac{1}{n_\ell-1} \sum_{i=0}^{n_\ell-1} (Y_\ell(\mathbf{x}_{\ell,i}) - \hat{\mu}_\ell)^2$ for $\ell \in \mathcal{L}$ \triangleright estimate of V_ℓ

$\mathcal{L}_{\text{feasible}} \leftarrow \{\ell \in \{1, \dots, L\} : \sum_{\ell'=1}^L C_{\ell'} n_{\ell'} + C_\ell n_\ell \leq N\}$

if $\mathcal{L}_{\text{feasible}} = \emptyset$ **then** break \triangleright exit loop before going over-budget

$\hat{\ell} \leftarrow \operatorname{argmax}_{\ell \in \mathcal{L}_{\text{feasible}}} \frac{\sigma_\ell^2}{n_\ell C_\ell}$ \triangleright chose the level of maximum utility

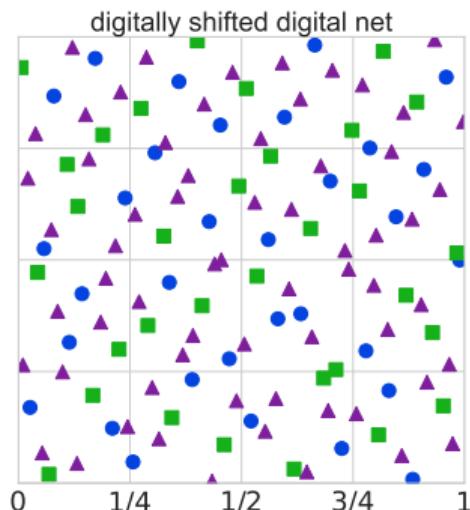
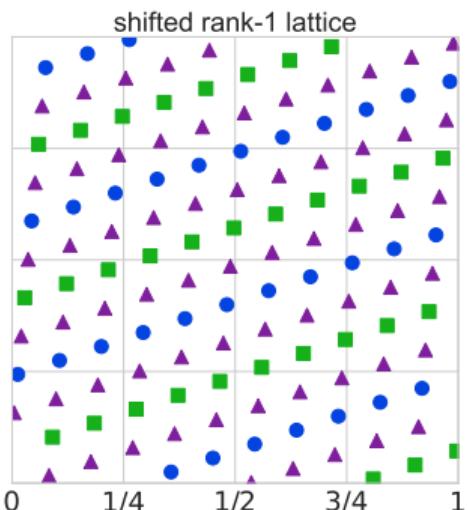
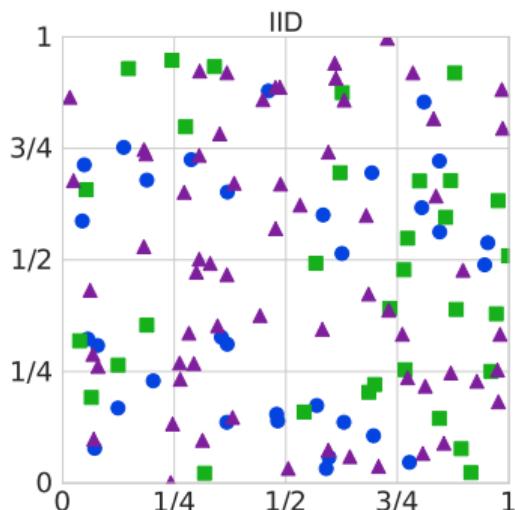
$\mathcal{L} \leftarrow \{\hat{\ell}\}$ and $n_{\hat{\ell}}^* \leftarrow n_{\hat{\ell}}^{\text{next}}$ and $n_{\hat{\ell}}^{\text{next}} \leftarrow 2n_{\hat{\ell}}^*$ \triangleright double samples on chosen level

$\hat{\nu} \leftarrow \sum_{\ell=1}^L \hat{\mu}_\ell$ and $\sigma^2 \leftarrow \sum_{\ell=1}^L \hat{\sigma}_\ell^2 / n_\ell$ \triangleright estimates of ν and $\mathbb{V}[\hat{\nu}]$ respectively

return $\hat{\nu}, \sigma, \{n_\ell\}_{\ell=1}^L$ \triangleright the estimate, its standard error, and samples per level

Quasi-Monte Carlo (QMC) Methods

QMC uses highly-uniform low-discrepancy (LD) sequences instead of IID points to accelerate Monte Carlo



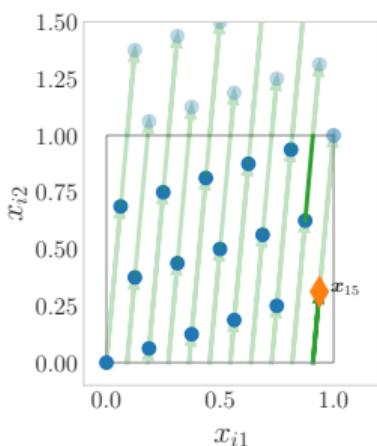
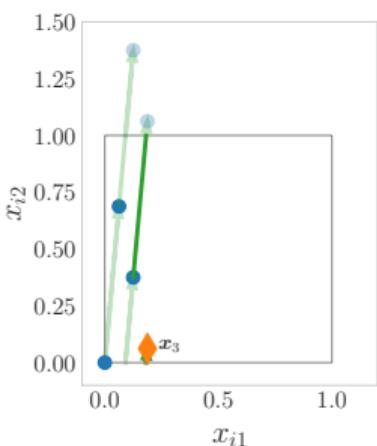
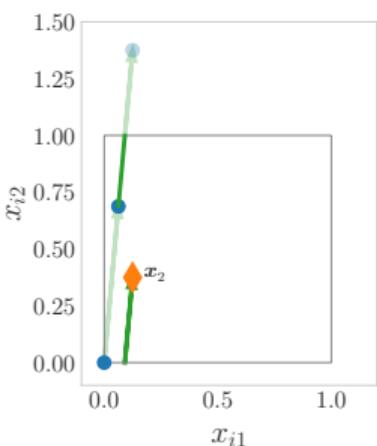
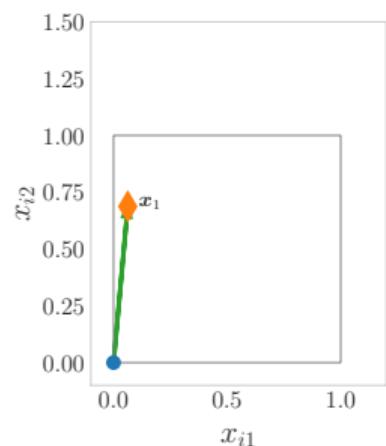
IID Monte Carlo has error like $\mathcal{O}(1/\sqrt{n})$. IID points have gaps and clusters.

QMC has errors like $\mathcal{O}(1/n)$. LD lattices and digital nets evenly cover the unit cube.

Unrandomized LD Lattice Construction for QMC

Generating vector $\mathbf{g} \in \mathbb{N}^d$ defines unrandomized lattice $(\mathbf{z}_i)_{i=0}^{n-1} \subset [0,1)^d$ with

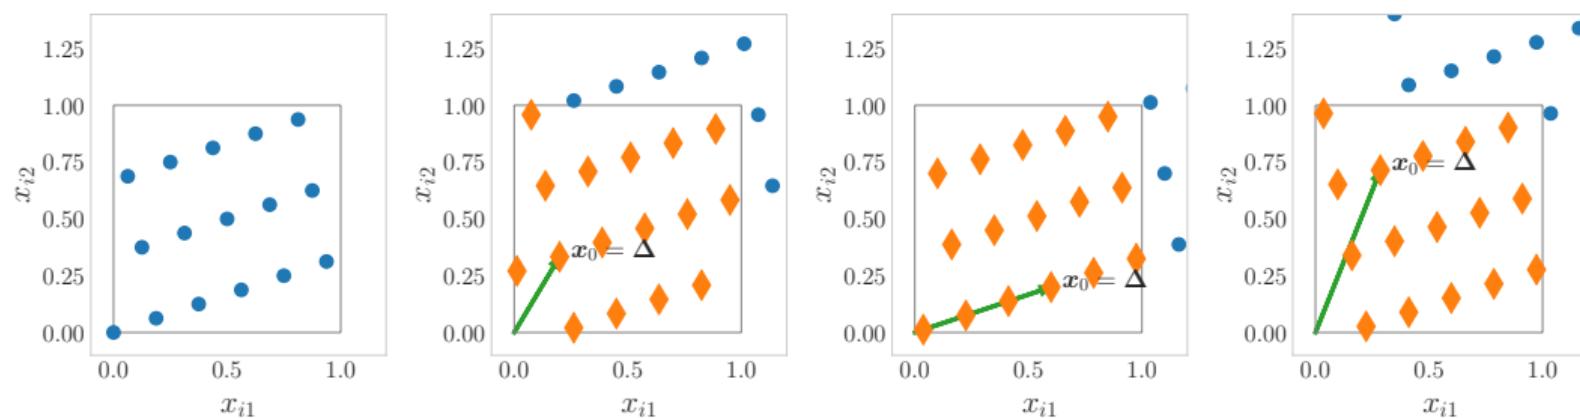
$$\mathbf{z}_i = (i\mathbf{g}/n) \mod \mathbf{1}, \quad i = 0, \dots, n-1.$$



Randomized LD Lattice Construction for QMC

A single random shift $\Delta \in [0, 1)^d$ defines the randomly shifted lattice $(\mathbf{x}_i)_{i=0}^{n-1} \subset [0, 1)^d$

$$\mathbf{x}_i = \mathbf{z}_i \oplus \Delta := (\mathbf{z}_i + \Delta) \mod 1, \quad i = 0, \dots, n-1.$$



$R > 1$ independently randomized LD sequences gives R independent QMC estimates [5]

Multilevel QMC (MLQMC) with Replications [6]

Existing MLQMC methods use $R > 1$ independently randomized (replications) of an LD sequence

$$\nu = \mathbb{E}[Q_L(\mathbf{X})] = \sum_{\ell=1}^L \mathbb{E}[Q_\ell(\mathbf{X}) - Q_{\ell-1}(\mathbf{X})] = \sum_{\ell=1}^L \mathbb{E}[Y_\ell(\mathbf{X})] = \sum_{\ell=1}^L \mu_\ell, \quad \mathbf{X} \sim \mathcal{U}[0,1]^d.$$

$R > 1$ independent shifts on each level gives $(\mathbf{x}_{\ell,r,i})_{i=0}^{n_\ell-1}$ for $1 \leq \ell \leq L$ and $1 \leq r \leq R$:

$$\mathbf{x}_{\ell,r,i} = \mathbf{z}_i \oplus \boldsymbol{\Delta}_{\ell,r}, \quad \boldsymbol{\Delta}_{1,1}, \dots, \boldsymbol{\Delta}_{1,R}, \dots, \boldsymbol{\Delta}_{L,1}, \dots, \boldsymbol{\Delta}_{L,R} \stackrel{\text{IID}}{\sim} \mathcal{U}[0,1]^d.$$

Approximate ν by the resulting LR independent sample means $\tilde{\mu}_{\ell,r}$,

$$\hat{\nu} = \sum_{\ell=1}^L \hat{\mu}_\ell \quad \text{where} \quad \hat{\mu}_\ell = \frac{1}{R} \sum_{r=1}^R \tilde{\mu}_{\ell,r} \quad \text{and} \quad \tilde{\mu}_{\ell,r} = \frac{1}{n_\ell} \sum_{i=0}^{n_\ell-1} Y_\ell(\mathbf{x}_{\ell,r,i}).$$

$$\therefore \mathbb{V}[\hat{\nu}] = \sum_{\ell=1}^L \tilde{V}_{\ell,n_\ell}/R \text{ for } \tilde{V}_{\ell,n_\ell} := \mathbb{V}\left[\frac{1}{n_\ell} \sum_{i=0}^{n_\ell-1} Y_\ell(\mathbf{z} \oplus \boldsymbol{\Delta})\right] \xrightarrow{n_\ell \rightarrow \infty} 0 \text{ for } \boldsymbol{\Delta} \sim \mathcal{U}[0,1]^d.$$

Algorithm 2 RQMC: MLQMC with Replications

Require: $N, C_\ell, n_\ell^{\text{next}}, R$ ▷ budget, costs, initial samples, replications

Generate $\Delta_{1,1}, \dots, \Delta_{1,R}, \dots, \Delta_{L,1}, \dots, \Delta_{L,r} \stackrel{\text{IID}}{\sim} \mathcal{U}[0,1]^d$

$n_\ell \leftarrow 0$ for $\ell \in \{1, \dots, L\}$ and $\mathcal{L} \leftarrow \{1, \dots, L\}$

while true **do**

 Generate $x_{\ell,r,i} = z_i \oplus \Delta_{\ell,r}$ for $\ell \in \mathcal{L}$ and $1 \leq r \leq R$ and $n_\ell \leq i < n_\ell^{\text{next}}$

 Evaluate $Y_\ell(x_{\ell,r,i})$ for $\ell \in \mathcal{L}$ and $1 \leq r \leq R$ and $n_\ell \leq i < n_\ell^{\text{next}}$

$\tilde{\mu}_{\ell,r} \leftarrow \frac{1}{n_\ell} \sum_{i=0}^{n_\ell-1} Y_\ell(x_{\ell,r,i})$ for $\ell \in \mathcal{L}$ and $1 \leq r \leq R$ ▷ per-randomization estimate of μ_ℓ

$\hat{\mu}_\ell \leftarrow \frac{1}{R} \sum_{r=1}^R \tilde{\mu}_{\ell,r}$ for $\ell \in \mathcal{L}$ ▷ aggregate estimate of mean μ_ℓ

$\tilde{\sigma}_\ell^2 \leftarrow \frac{1}{R-1} \sum_{r=1}^R (\tilde{\mu}_{\ell,r} - \hat{\mu}_\ell)^2$ for $\ell \in \mathcal{L}$ ▷ estimate of variance \tilde{V}_{ℓ,n_ℓ}

$\mathcal{L}_{\text{feasible}} \leftarrow \{\ell \in \{1, \dots, L\} : \sum_{\ell'=1}^L R C_{\ell'} n_{\ell'} + R C_\ell n_\ell \leq N\}$

if $\mathcal{L}_{\text{feasible}} = \emptyset$ **then break** ▷ exit loop before going over budget

$\ell^* \leftarrow \operatorname{argmax}_{\ell \in \mathcal{L}_{\text{feasible}}} \frac{\tilde{\sigma}_\ell^2}{R n_\ell C_\ell}$ ▷ choose the level of maximum utility

$\mathcal{L} \leftarrow \{\ell^*\}$ and $n_{\ell^*} \leftarrow n_{\ell^*}^{\text{next}}$ and $n_{\ell^*}^{\text{next}} \leftarrow 2n_{\ell^*}$ ▷ double samples on chosen level

$\hat{\nu} \leftarrow \sum_{\ell=1}^L \hat{\mu}_\ell$ and $\tilde{\sigma}^2 \leftarrow \sum_{\ell=1}^L \tilde{\sigma}_\ell^2 / R$ ▷ estimate for ν and $\mathbb{V}[\hat{\nu}]$ respectively

return $\hat{\nu}, \tilde{\sigma}, \{R n_\ell\}_{\ell=1}^L$ ▷ the estimate, its standard error, and samples per level

QMC with Multiple Replications is Inefficient

Bayesian cubature enables QMC error estimation without inefficient replications

- Most QMC methods use replications only for error estimation
- Fewer replications R gives a more accurate QMC estimate $\hat{\mu}_\ell$, but a worse estimate for the predicted error $\tilde{\sigma}_\ell$
- Replications R trades off accuracy of the approximation for accuracy of the error estimate
- MLQMC needs both accurate approximations and estimates

Can we avoid replications while still getting an error estimate?

Yes! Using Bayesian cubature with LD sampling locations

- A single LD sequence is used to construct a Gaussian process (GP).
- Error estimation using GPs posterior variance.
- Bonus: GPs with matching kernels and LD points cost only $\mathcal{O}(n \log n)$ to fit

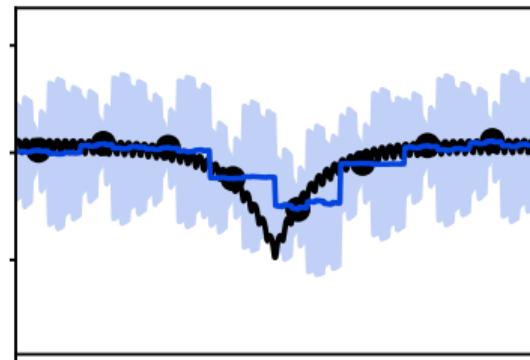
Gaussian Process Regression [7]

Gaussian process $Y \sim \text{GP}(\tau, K)$ has posterior mean and covariance

$$\mathbb{E}[Y(\mathbf{x})|\mathbf{X}, \mathbf{Y}] = \tau + \mathbf{K}^\top(\mathbf{x}) \mathbf{K}^{-1}(\mathbf{Y} - \tau \mathbf{1}) \quad \text{and}$$

$$\text{Cov}[Y(\mathbf{x}), Y(\mathbf{x}')|\mathbf{X}] = K(\mathbf{x}, \mathbf{x}') - \mathbf{K}^\top(\mathbf{x}) \mathbf{K}^{-1} \mathbf{K}(\mathbf{x}')$$

- Constant prior mean $\tau \in \mathbb{R}$
- Nodes $\mathbf{X} := (\mathbf{x}_i)_{i=0}^{n-1}$
- Evaluations $\mathbf{Y} := (Y(\mathbf{x}_i))_{i=0}^{n-1} \in \mathbb{R}^n$
- Kernel vector $\mathbf{K}(\mathbf{x}) = (K(\mathbf{x}, \mathbf{x}_i))_{i=0}^{n-1} \in \mathbb{R}^n$
- SPD Gram matrix $\mathbf{K} = (K(\mathbf{x}_i, \mathbf{x}_{i'}))_{i,i'=0}^{n-1} \in \mathbb{R}^{n \times n}$



Kernel hyperparameters $\boldsymbol{\theta}$ often optimized with NMLL loss

$$L(\boldsymbol{\theta}) = \log |\mathbf{K}| + (\mathbf{Y} - \tau \mathbf{1})^\top \mathbf{K}^{-1} (\mathbf{Y} - \tau \mathbf{1})$$

GP regression typically requires $\mathcal{O}(n^2)$ storage and $\mathcal{O}(n^3 + n^2d)$ computations.

Fast GPs Pairing LD Points with Special Kernels [8, 9]

Accelerate GP regression by inducing nice structure into the Gram matrix K

1. LD Lattices + Shift Invariant (SI) Kernels

- Give circulant Gram matrices $K = (K(\mathbf{x}_i, \mathbf{x}_{i'}))_{i,i'=0}^{n-1}$
- ∴ Eigendecomp $K = V \Lambda \bar{V}$ where \bar{V} is the DFT matrix → FFT in $\mathcal{O}(n \log n)$

2. LD Digital Nets + Digitally Shift Invariant (DSI) Kernels

- Give Recursive Symmetric Block Toeplitz (RSBT) Gram matrices K
- ∴ Eigendecomp $K = V \Lambda \bar{V}$ where \bar{V} is the Hadamard matrix → FWHT in $\mathcal{O}(n \log n)$

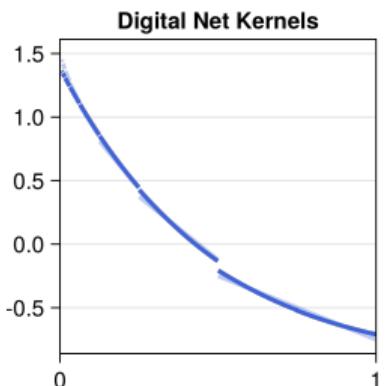
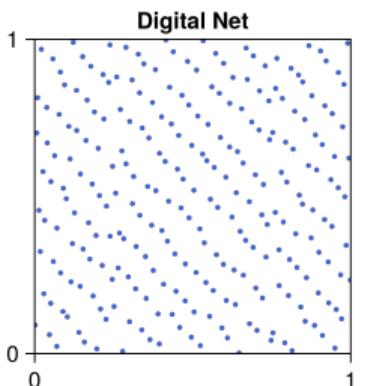
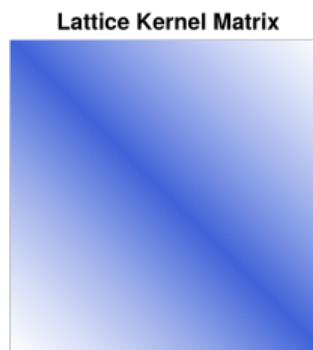
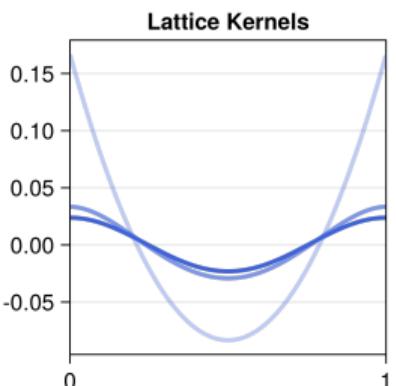
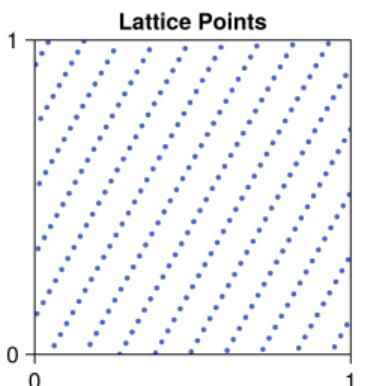
Let $v_1 = \mathbf{1}/\sqrt{n}$ and k_1 be the first columns of \bar{V} and K respectively:

$$\lambda := \Lambda \mathbf{1} = \sqrt{n} \Lambda v_1 = \sqrt{n} \bar{V} V \Lambda v_1 = \sqrt{n} \bar{V} k_1$$

- Ka , $K^{-1}a$, and $|K|$ can all be computed in $\mathcal{O}(n \log n)$ computations
- Only requires evaluating and storing the first column of K

Fast GP regression only requires $\mathcal{O}(n)$ storage and $\mathcal{O}(n \log n + nd)$ computations.

Lattices + SI $K =$ Circulant K and Digital Nets + DSI $K =$ RSBT K



Fast Bayesian Cubature [10, 11, 12]

$$\mu = \mathbb{E}_{\mathbf{X}}[Y(\mathbf{X})] \sim \mathcal{N}(\hat{\mu}, \hat{V}_n), \quad \mathbf{X} \sim \mathcal{U}[0, 1]^d$$

with respective posterior cubature mean and variance

$$\hat{\mu} := \mathbb{E}[\mu | \mathbf{X}, \mathbf{Y}] = \tau + \left(\int \mathbf{K}(\mathbf{x}) d\mathbf{x} \right)^\top \mathbf{K}^{-1} (\mathbf{Y} - \tau \mathbf{1}) \quad \text{and}$$

$$\hat{V}_n := \mathbb{V}[\mu | \mathbf{X}] = \iint K(\mathbf{x}, \mathbf{x}') d\mathbf{x} d\mathbf{x}' - \left(\int \mathbf{K}(\mathbf{x}) d\mathbf{x} \right)^\top \mathbf{K}^{-1} \left(\int \mathbf{K}(\mathbf{x}) d\mathbf{x} \right).$$

Certain fast GP constructions make the GP prediction $\hat{\mu}$ equal the QMC sample mean:

$$\hat{\mu} = \tau = \frac{1}{n} \sum_{i=0}^{n-1} Y(\mathbf{x}_i)$$

Fast Bayesian MLQMC without Replications

$$\mu_\ell = \mathbb{E}_{\mathbf{X}}[Y_\ell(\mathbf{X})] \sim \mathcal{N}(\hat{\mu}_\ell, \hat{V}_{\ell, n_\ell}), \quad \mathbf{X} \sim \mathcal{U}[0, 1]^d$$

Posterior cubature mean

$$\hat{\nu} := \mathbb{E}[\nu | (\mathbf{X}_\ell, \mathbf{Y}_\ell)_{\ell=1}^L] = \sum_{\ell=1}^L \hat{\mu}_\ell \quad \text{where} \quad \hat{\mu}_\ell := \mathbb{E}[\mu_\ell | \mathbf{X}_\ell, \mathbf{Y}_\ell] = \frac{1}{n_\ell} \sum_{i=0}^{n_\ell-1} Y_\ell(\mathbf{x}_{\ell,i})$$

and computable posterior cubature variance (assuming independent GPs)

$$\hat{V}_{n_1, \dots, n_L} := \mathbb{V}[\nu | (\mathbf{X}_\ell)_{\ell=1}^L] = \sum_{\ell=1}^L \hat{V}_{\ell, n_\ell} \quad \text{where} \quad \hat{V}_{\ell, n_\ell} \text{ only depends on } \mathbf{X}_\ell \quad (1)$$

Use projected variance, $\hat{V}_{\ell, \hat{n}_\ell}$ with $\hat{n}_\ell \geq n_\ell$, to decide on which level to double samples

Algorithm 3 BQMC: Fast Bayesian MLQMC Without Replications (Independent GPs)

Require: $N, C_\ell, n_\ell^{\text{next}}, \theta_\ell$ \triangleright budget, costs, initial samples, initial kernel parameters

$\Delta_1, \dots, \Delta_L \stackrel{\text{IID}}{\sim} \mathcal{U}[0,1]^d$ \triangleright random shifts, one per level

$n_\ell \leftarrow 0$ for $\ell \in \{1, \dots, L\}$ and $\mathcal{L} \leftarrow \{1, \dots, L\}$

while true **do**

 Generate $x_{\ell,i} = z_i \oplus \Delta_\ell$ for $\ell \in \mathcal{L}$ and $n_\ell \leq i < n_\ell^{\text{next}}$

 Evaluate $Y_\ell(x_{\ell,i})$ for $\ell \in \mathcal{L}$ and $n_\ell \leq i < n_\ell^{\text{next}}$

 Update θ_ℓ to optimize the NMLL for $\ell \in \mathcal{L}$

$\hat{\mu}_\ell \leftarrow \frac{1}{n_\ell} \sum_{i=0}^{n_\ell-1} Y_\ell(x_{\ell,i})$ for $\ell \in \mathcal{L}$ \triangleright per-level posterior cubature mean

 Compute \hat{V}_{ℓ,n_ℓ} for $\ell \in \mathcal{L}$ \triangleright per-level posterior cubature variance

$\mathcal{L}_{\text{feasible}} \leftarrow \{\ell \in \{1, \dots, L\} : \sum_{\ell'=1}^L C_{\ell'} n_{\ell'} + C_\ell n_\ell \leq N\}$

if $\mathcal{L}_{\text{feasible}} = \emptyset$ **then** break \triangleright exit loop before going over budget

$\ell^\star \leftarrow \text{level_select_BQMC}(\mathcal{L}_{\text{feasible}}, (C_\ell, \theta_\ell, n_\ell)_{\ell=1}^L)$ \triangleright choose level with Algorithm 4

$\mathcal{L} \leftarrow \{\ell^\star\}$ and $n_{\ell^\star} \leftarrow n_{\ell^\star}^{\text{next}}$ and $n_{\ell^\star}^{\text{next}} \leftarrow 2n_{\ell^\star}$ \triangleright double samples on chosen level

$\hat{\nu} \leftarrow \sum_{\ell=1}^L \hat{\mu}_\ell$ and $\hat{V}_{n_1, \dots, n_L} \leftarrow \sum_{\ell=1}^L \hat{V}_{\ell,n_\ell}$ \triangleright posterior cubature mean and variance respectively

return $\hat{\nu}, \sqrt{\hat{V}_{n_1, \dots, n_L}}, \{n_\ell\}_{\ell=1}^L$ \triangleright the estimate, its standard error, and samples per level

Numerical Experiments: Single-Level Functions in $d = 32$ Dimensions

Ridge functions $Y(\mathbf{x}) = g(u(\mathbf{x}))$ use $u(\mathbf{x}) = \sum_{j=1}^d c_j \Phi^{-1}(x_j)$ with weights $c_j = 2^{-j}/\sqrt{\sum_{j'=1}^d 2^{-2j'}}$ and Φ the normal CDF.

1. **sumxex** $Y(\mathbf{x}) = -d + \sum_{j=1}^d x_j \exp(x_j)$. Smooth and additive. Easy for QMC.
2. **ridge PL sparse** $Y(\mathbf{x}) = g(u(\mathbf{x}))$ with $g(u) = \max(u - 1, 0) - \phi(1) + \Phi(-1)$ and ϕ the density of $\mathcal{N}(0, 1)$. Continuous piecewise linear (PL) with a kink.
3. **ridge JSU sparse** $Y(\mathbf{x}) = g(u(\mathbf{x}))$ with $g(u) = -\eta + F^{-1}(\Phi(v))$ where F is the CDF of a Johnson's SU distribution [13]. Heavy tailed.
4. **Genz corner-peak 2** $Y(\mathbf{x}) = \left(1 + \sum_{j=1}^d c_j x_j\right)^{-(d+1)}$ with coefficients of the second kind $c_j = j^{-2}/(4 \sum_{j'=1}^d (j')^{-2})$. A Genz function [14] with moderate anisotropy and effective dimension.

Numerical Experiments: Multilevel Problems

Decay of the means μ_ℓ and standard deviations $\sqrt{V_\ell}$

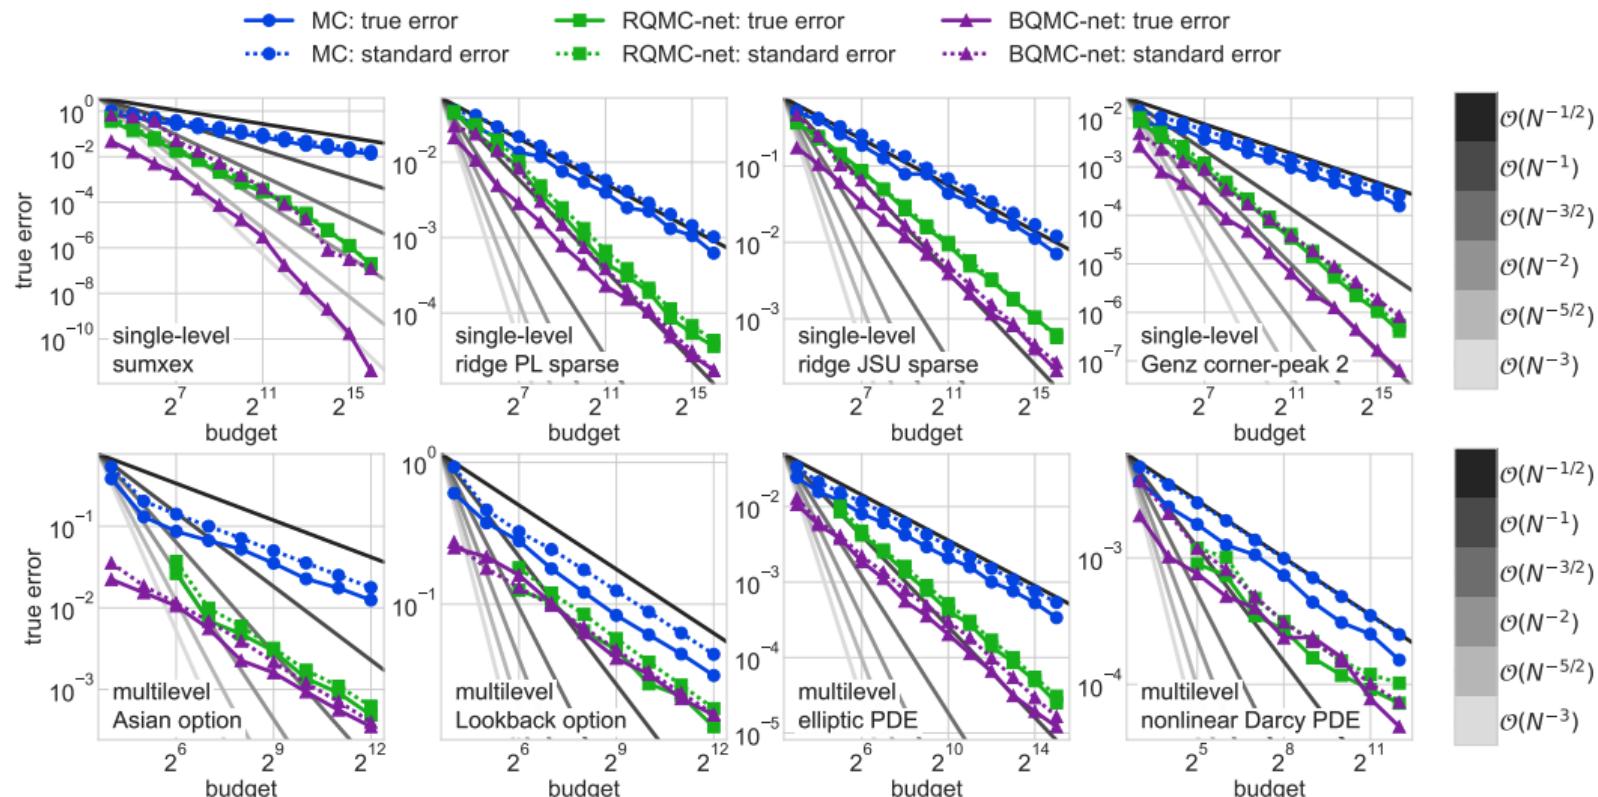
mean $\mu_\ell = \mathbb{E}[Y_\ell(\mathbf{X})] = \mathbb{E}[Q_\ell(\mathbf{X}) - Q_{\ell-1}(\mathbf{X})]$ with $\mathbf{X} \sim \mathcal{U}[0, 1]^d$

problem/ ℓ	1	2	3	4	5	6	7	8
Asian option	6.3e0	-3.0e-1	-1.5e-1	-7.2e-2	-3.7e-2	-1.8e-2	-9.3e-3	-4.8e-3
Lookback option	1.3e1	1.4e0	8.9e-1	5.8e-1	4.1e-1	2.8e-1	1.8e-1	1.3e-1
elliptic PDE	1.6e-1	-1.1e-2	2.5e-3	1.5e-3				
nonlinear Darcy PDE	4.1e-2	8.5e-4	3.3e-4					

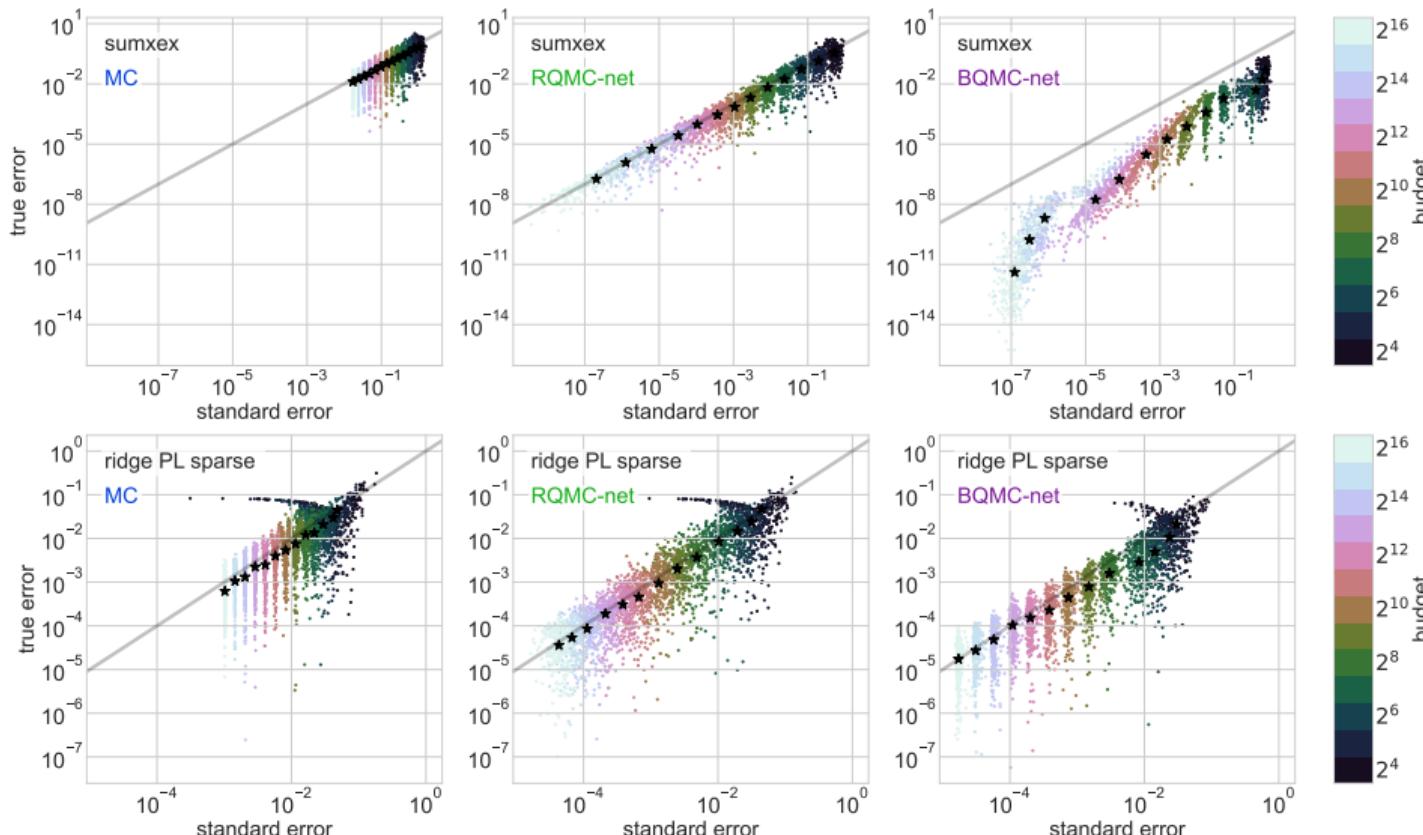
standard deviation $\sqrt{V_\ell} = \sqrt{\mathbb{V}[Y_\ell(\mathbf{X})]} = \sqrt{\mathbb{V}[Q_\ell(\mathbf{X}) - Q_{\ell-1}(\mathbf{X})]}$ with $\mathbf{X} \sim \mathcal{U}[0, 1]^d$

problem/ ℓ	1	2	3	4	5	6	7	8
Asian option	8.7e0	6.8e-1	3.5e-1	1.7e-1	8.4e-2	4.3e-2	2.1e-2	1.1e-2
Lookback option	1.3e1	2.1e0	1.4e0	9.0e-1	6.3e-1	4.2e-1	2.9e-1	2.0e-1
elliptic PDE	1.4e-1	6.2e-2	1.0e-2	3.5e-3				
nonlinear Darcy PDE	7.8e-2	6.7e-3	1.9e-3					

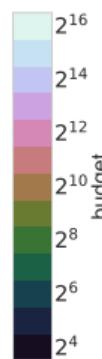
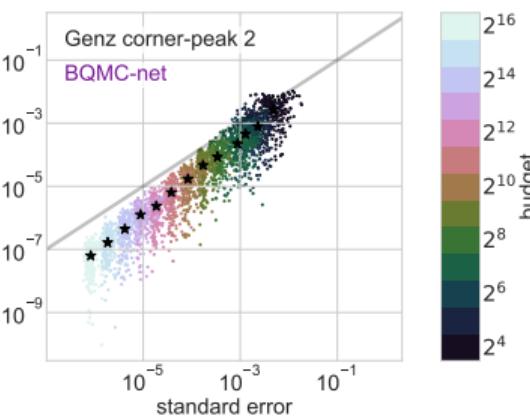
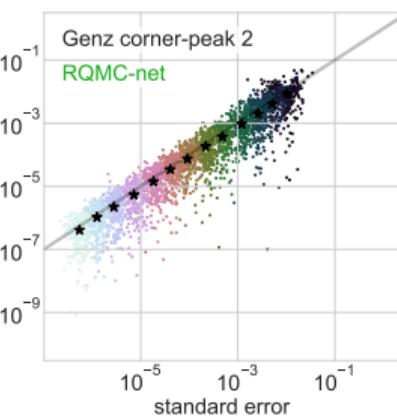
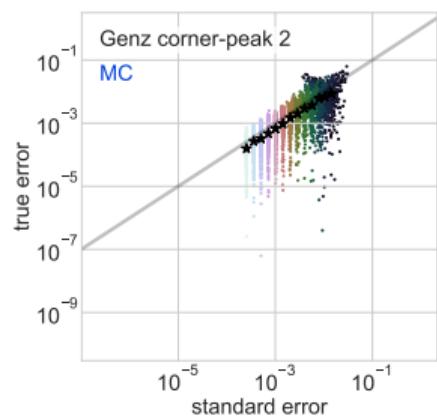
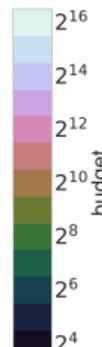
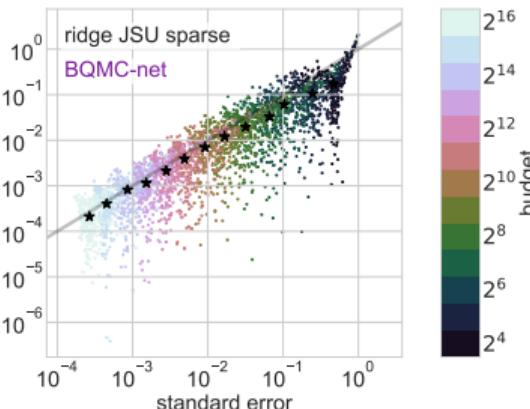
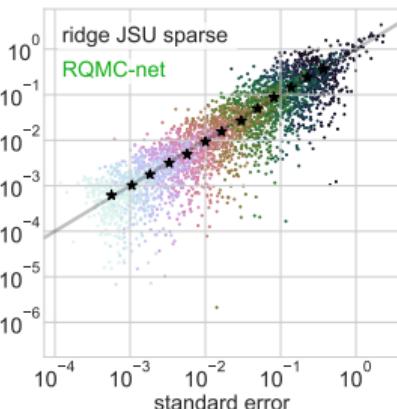
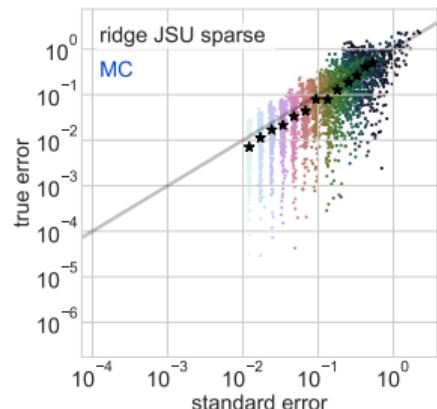
Convergence with Increasing Budgets: LD Digital Nets



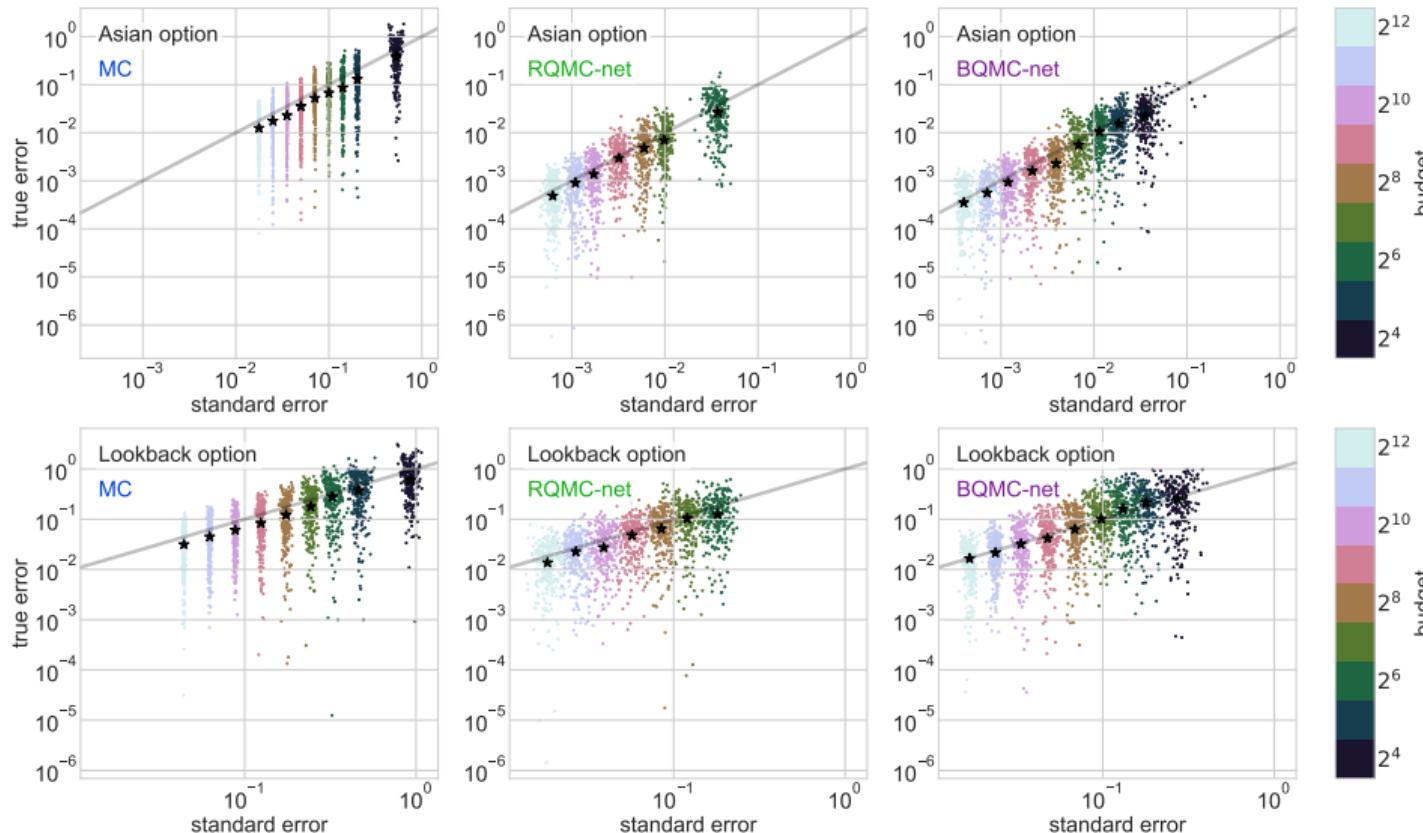
Single-Level Error Estimate Accuracy: LD Digital Nets



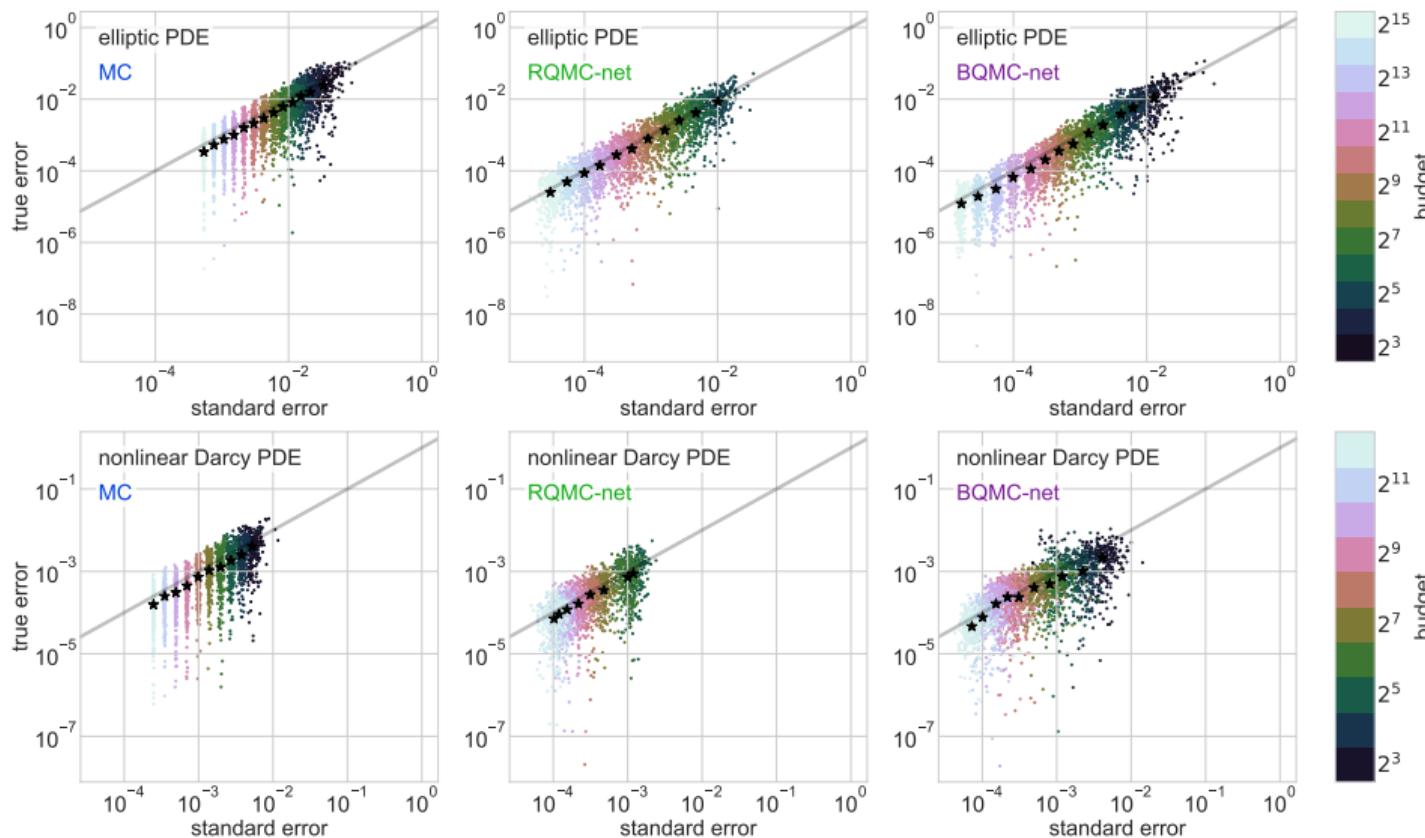
Single-Level Error Estimate Accuracy: LD Digital Nets



Multilevel Error Estimate Accuracy: LD Digital Nets

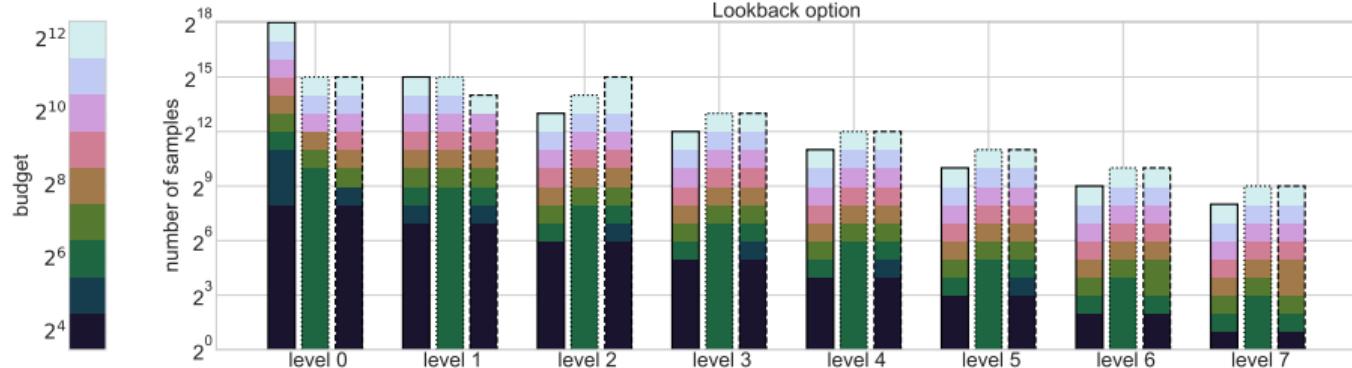
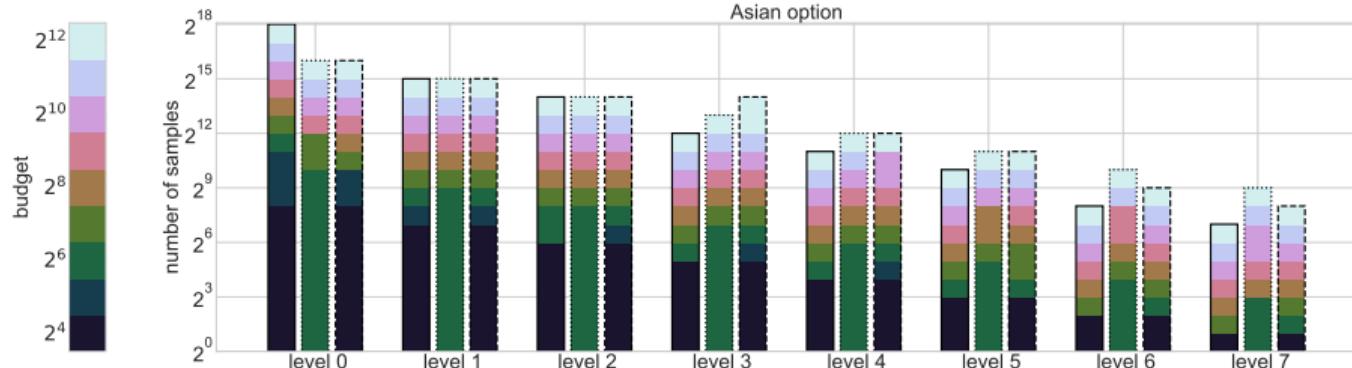


Multilevel Error Estimate Accuracy: LD Digital Nets

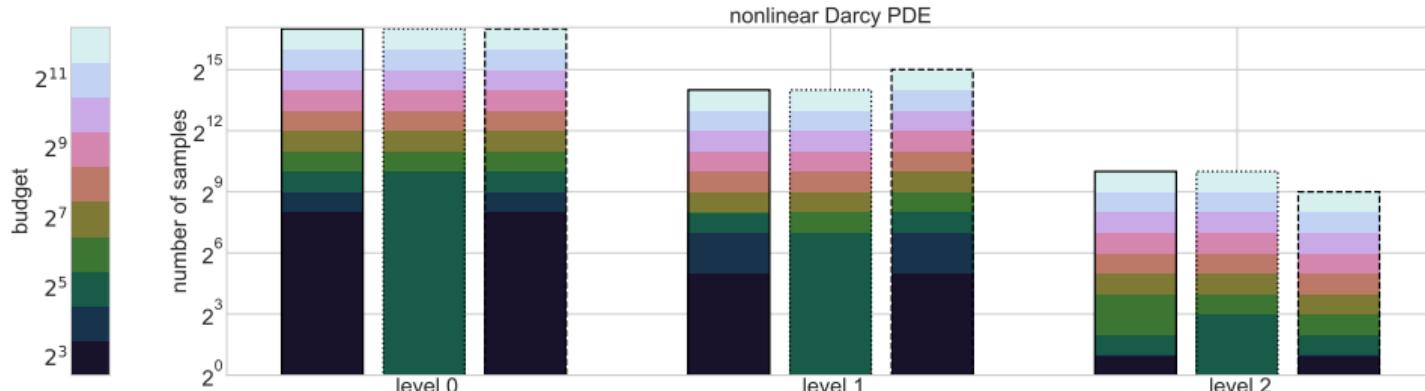
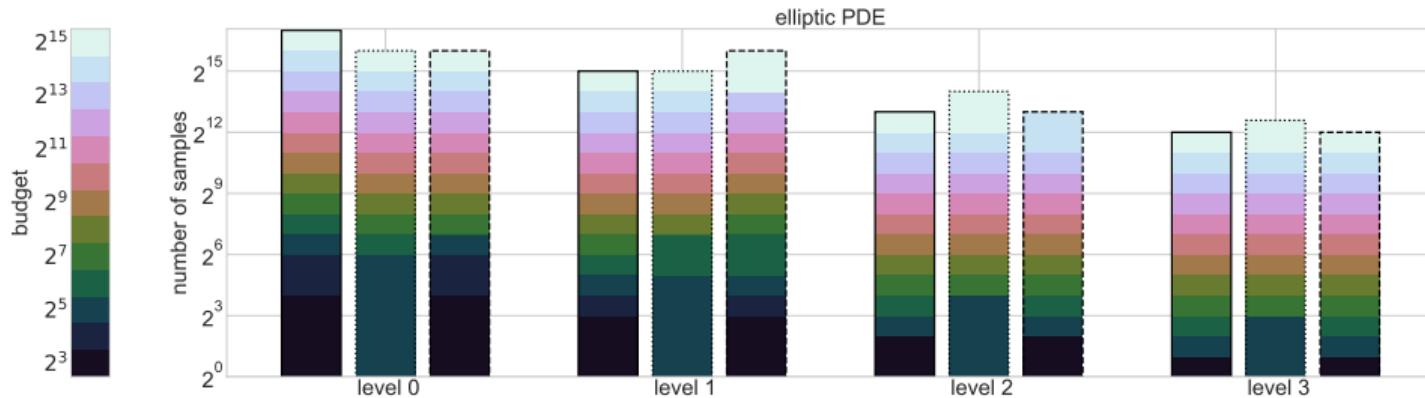


Multilevel Sample Allocation: LD Digital Nets

— MC - - - RQMC-net - - - BQMC-net



Multilevel Sample Allocation: LD Digital Nets



Summary and Software

[1] <https://arxiv.org/abs/2510.24604>

Algorithms

1. Multilevel Monte Carlo with IID points ([MLMC](#))
2. Multilevel Quasi-Monte Carlo with Replications ([RMLQMC](#))
3. Fast Bayesian Multilevel Quasi-Monte Carlo without Replications ([BMLQMC](#))

Python Software

1. [FastGPs](#): fast $\mathcal{O}(n \log n)$ GP regression using matching kernels and LD points
<https://alegresor.github.io/fastgps/>
2. [QMCPy](#): QMC tools for LD points, transforms, stopping criteria, ... [15, 16, 17, 18]
<https://qmcsoftware.github.io/QMCSoftware/>

Thank you for listening!

References I

- [1] Aleksei G. Sorokin, Pieterjan Robbe, Gianluca Geraci, Michael S. Eldred, and Fred J. Hickernell. Fast Bayesian multilevel quasi-Monte Carlo. *ArXiv preprint*, abs/2510.24604, 2025. URL <https://arxiv.org/abs/2510.24604>.
- [2] Kenneth Levenberg. A method for the solution of certain non-linear problems in least squares. *Quarterly of Applied Mathematics*, 2(2):164–168, 1944.
- [3] Donald W. Marquardt. An algorithm for least-squares estimation of nonlinear parameters. *Journal of the Society for Industrial and Applied Mathematics*, 11(2):431–441, 1963.
- [4] Michael B. Giles. Multilevel Monte Carlo path simulation. *Operations research*, 56(3):607–617, 2008.
- [5] Pierre L’Ecuyer, Marvin K. Nakayama, Art B. Owen, and Bruno Tuffin. Confidence intervals for randomized quasi-Monte Carlo estimators. In *Proceedings of the Winter Simulation Conference 2023*, pages 445–456. IEEE, 2023.

References II

- [6] Michael B. Giles and Benjamin J. Waterhouse. Multilevel quasi-Monte Carlo path simulation. *Advanced Financial Modelling, Radon Series on Computational and Applied Mathematics*, 8:165–181, 2009.
- [7] Carl Edward Rasmussen and Christopher K. I. Williams. *Gaussian processes for machine learning*. MIT Press, Cambridge, Massachusetts, 2006. URL <http://www.gaussianprocess.org/gpml/>.
- [8] Xiaoyan Zeng, Peter Kritzer, and Fred J. Hickernell. Spline methods using integration lattices and digital nets. *Constructive Approximation*, 30:529–555, 2009.
- [9] Xiaoyan Zeng, King-Tai Leung, and Fred J. Hickernell. *Error analysis of splines for periodic problems using lattice designs*. Springer, 2006.

References III

- [10] Jagadeeswaran Rathinavel and Fred J. Hickernell. Fast automatic Bayesian cubature using lattice sampling. *Statistics and Computing*, 29(6):1215–1229, 2019. ISSN 1573-1375. doi: 10.1007/s11222-019-09895-9. URL <http://dx.doi.org/10.1007/s11222-019-09895-9>.
- [11] Jagadeeswaran Rathinavel and Fred J. Hickernell. Fast automatic Bayesian cubature using Sobol' sampling. In *Advances in Modeling and Simulation: Festschrift for Pierre L'Ecuyer*, pages 301–318. Springer, 2022.
- [12] Jagadeeswaran Rathinavel. *Fast automatic Bayesian cubature using matching kernels and designs*. PhD thesis, Illinois Institute of Technology, 2019.
- [13] Norman L. Johnson. Systems of frequency curves generated by methods of translation. *Biometrika*, 36(1/2):149–176, 1949.
- [14] Alan Genz. Comparison of methods for the computation of multivariate normal probabilities. *Computing Science and Statistics*, pages 400–400, 1993.

References IV

- [15] Sou-Cheng T. Choi, Fred J. Hickernell, Jagadeeswaran Rathinavel, Michael J. McCourt, and Aleksei G. Sorokin. Quasi-Monte Carlo software. In Alexander Keller, editor, *Monte Carlo and Quasi-Monte Carlo Methods 2020*, pages 23–47, Cham, 2022. Springer International Publishing. ISBN 978-3-030-98319-2.
- [16] Aleksei G. Sorokin and Jagadeeswaran Rathinavel. On bounding and approximating functions of multiple expectations using quasi-Monte Carlo. In Aicke Hinrichs, Peter Kritzer, and Friedrich Pillichshammer, editors, *Monte Carlo and Quasi-Monte Carlo Methods 2022*, pages 583–599. Springer, 2022.
- [17] Sou-Cheng T. Choi, Yuhang Ding, Fred J. Hickernell, Jagadeeswaran Rathinavel, and Aleksei G. Sorokin. Challenges in developing great quasi-Monte Carlo software. In Aicke Hinrichs, Peter Kritzer, and Friedrich Pillichshammer, editors, *Monte Carlo and Quasi-Monte Carlo Methods 2022*, pages 209–222. Springer, 2022.

References V

- [18] Fred J. Hickernell, Nathan Kirk, and Aleksei G. Sorokin. Quasi-Monte Carlo methods: what, why, and how? *ArXiv preprint*, abs/2502.03644, 2025. URL <https://arxiv.org/abs/2502.03644>.

Algorithm 4 level_select_BQMC: Level selection for fast Bayesian MLQMC

Require: $\mathcal{L}_{\text{feasible}} \subseteq \{1, \dots, L\}$ with $\tilde{L} := |\mathcal{L}_{\text{feasible}}| > 0$ elements \triangleright levels to consider

Require: $C_\ell, \theta_\ell, n_\ell$ \triangleright costs, GP hyperparameters, current samples

Set unique $l_1, \dots, l_{\tilde{L}} \in \mathcal{L}_{\text{feasible}}$ so that $n_{l_1} C_{l_1} \geq \dots \geq n_{l_{\tilde{L}}} C_{l_{\tilde{L}}}$
 \triangleright order feasible levels by non-increasing cost-of-doubling

$\ell \leftarrow l_1$ \triangleright initialize the selected level

for $k = 2, \dots, \tilde{L}$ **do**

$\ell' \leftarrow \ell_k$ \triangleright will compare levels ℓ and ℓ' with $n_\ell C_\ell \geq n_{\ell'} C_{\ell'}$

$\hat{n}_{\ell'} \leftarrow n_\ell C_\ell / C_{\ell'} + n_{\ell'}$ \triangleright gives equal costs for increasing sample sizes

$p \leftarrow \lfloor \log_2(\hat{n}_{\ell'}) \rfloor$ \triangleright implies $2^p \leq \hat{n}_{\ell'} < 2^{p+1}$

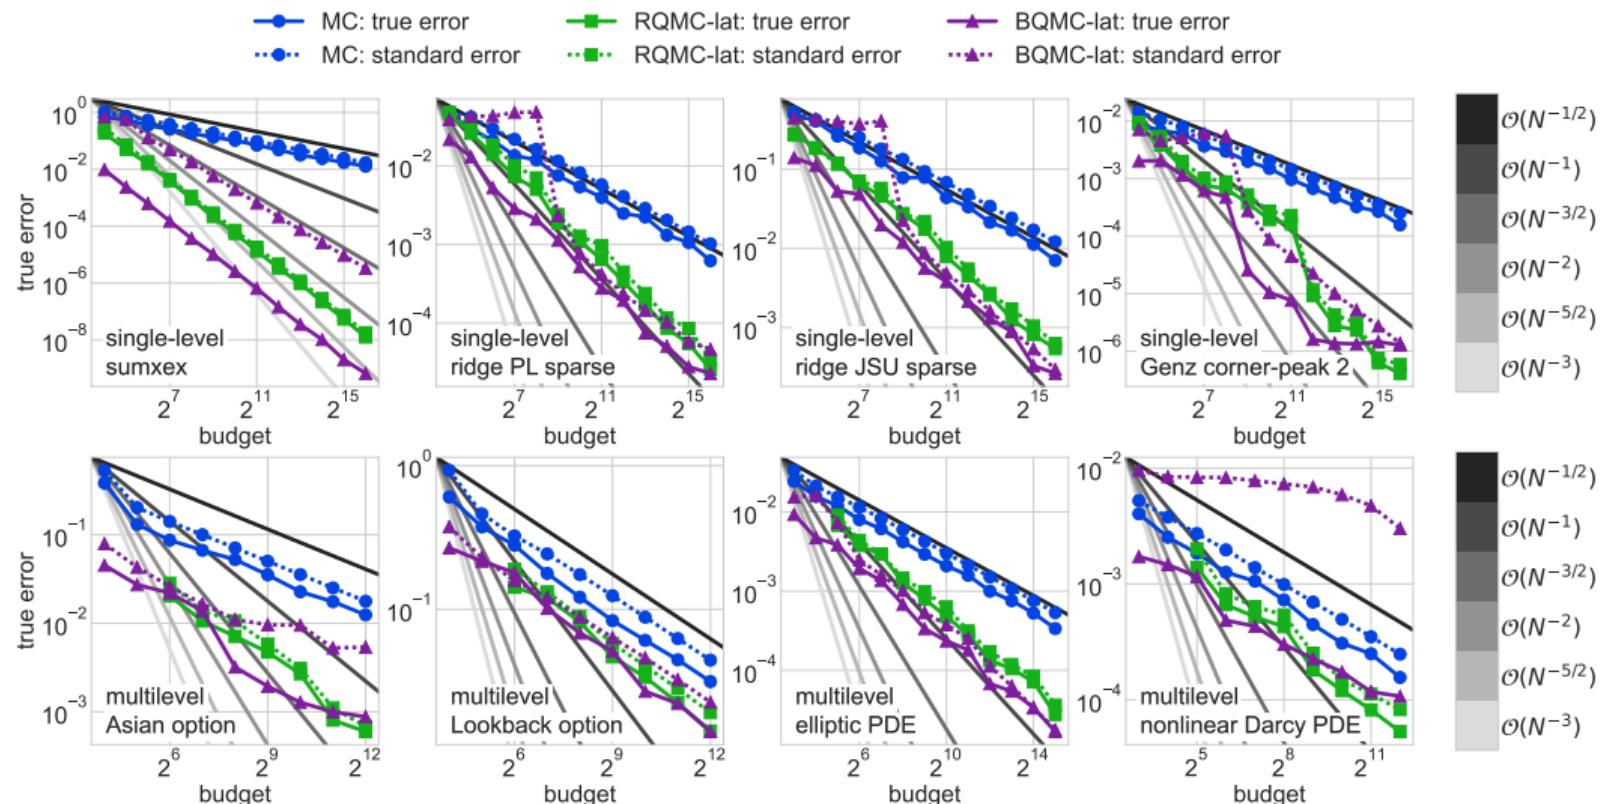
$\hat{V}_{\ell', \hat{n}_{\ell'}} \leftarrow (\hat{V}_{\ell', 2^p}^{p+1} / \hat{V}_{\ell', 2^{p+1}}^p) \hat{n}_{\ell'}^{\log_2(\hat{V}_{\ell', 2^{p+1}} / \hat{V}_{\ell', 2^p})}$ \triangleright log-log interpolation

if $\hat{V}_{\ell', n_{\ell'}} - \hat{V}_{\ell', \hat{n}_{\ell'}} \geq \hat{V}_{\ell, n_\ell} - \hat{V}_{\ell, 2n_\ell}$ **then** \triangleright a greater projected variance decrease

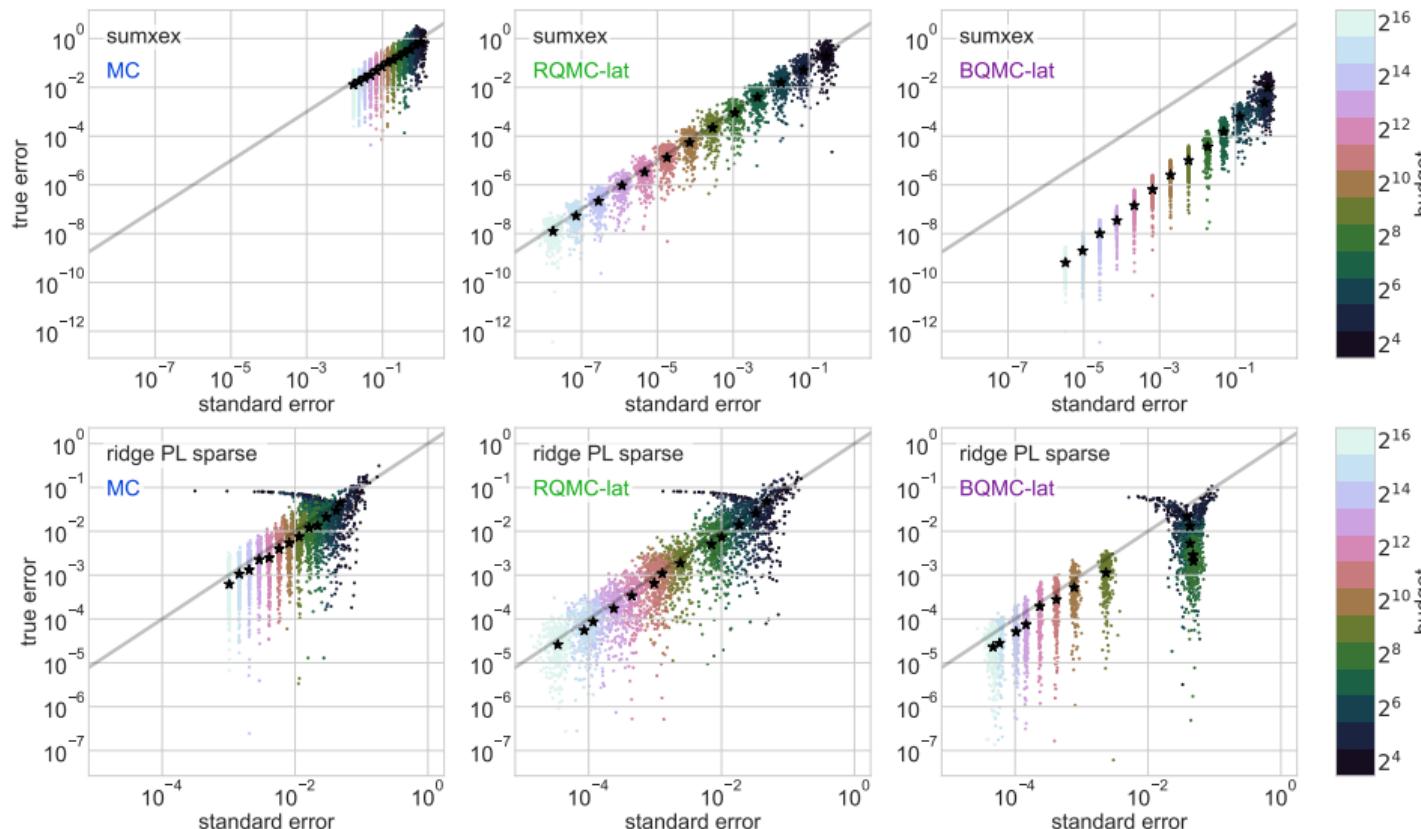
$\ell \leftarrow \ell'$ \triangleright update selected level

return ℓ \triangleright the selected level

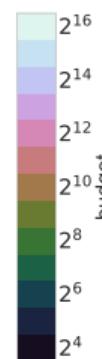
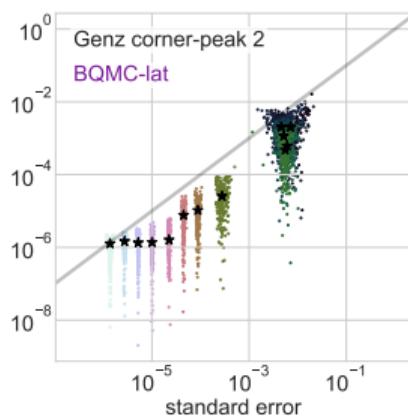
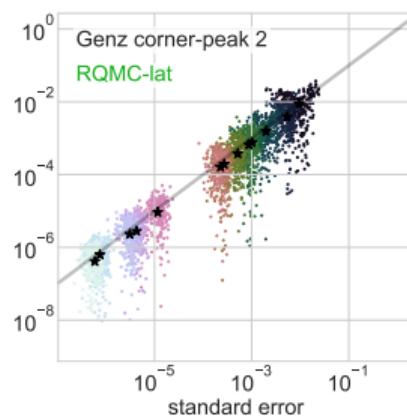
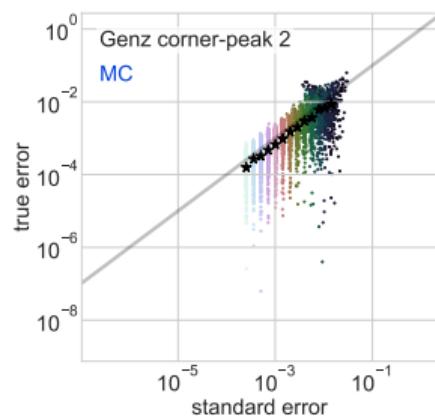
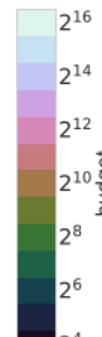
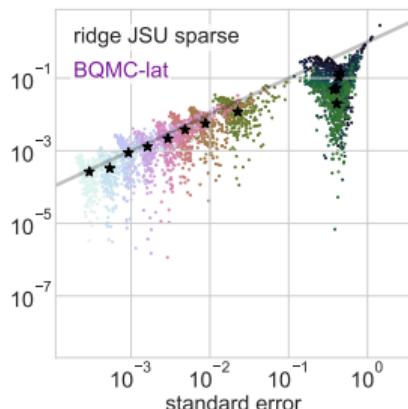
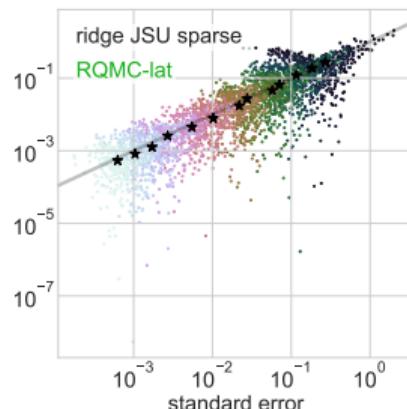
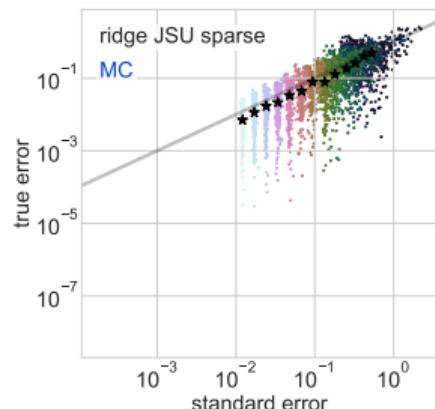
Convergence with Increasing Budgets: LD Lattices



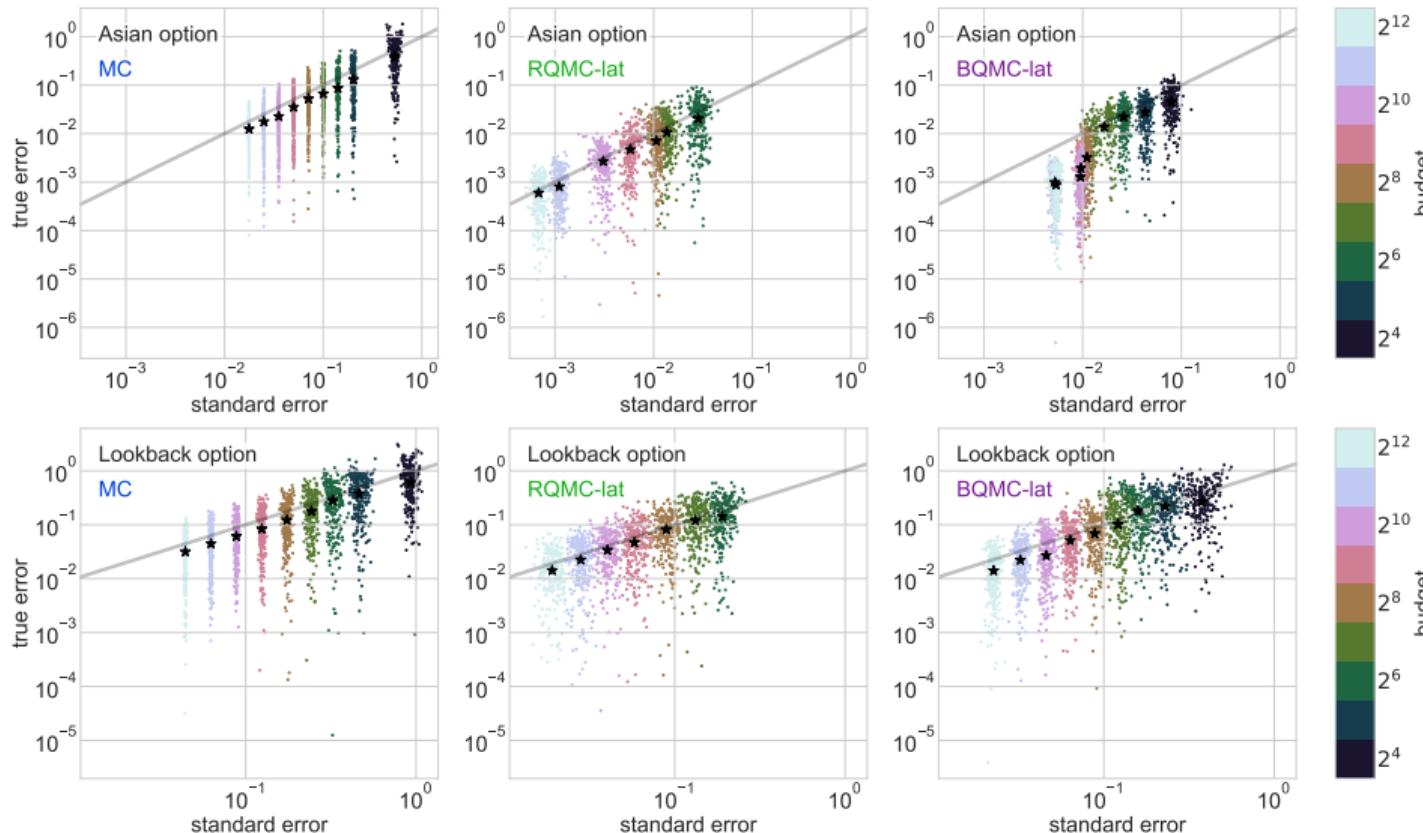
Single-Level Error Estimate Accuracy: LD Lattices



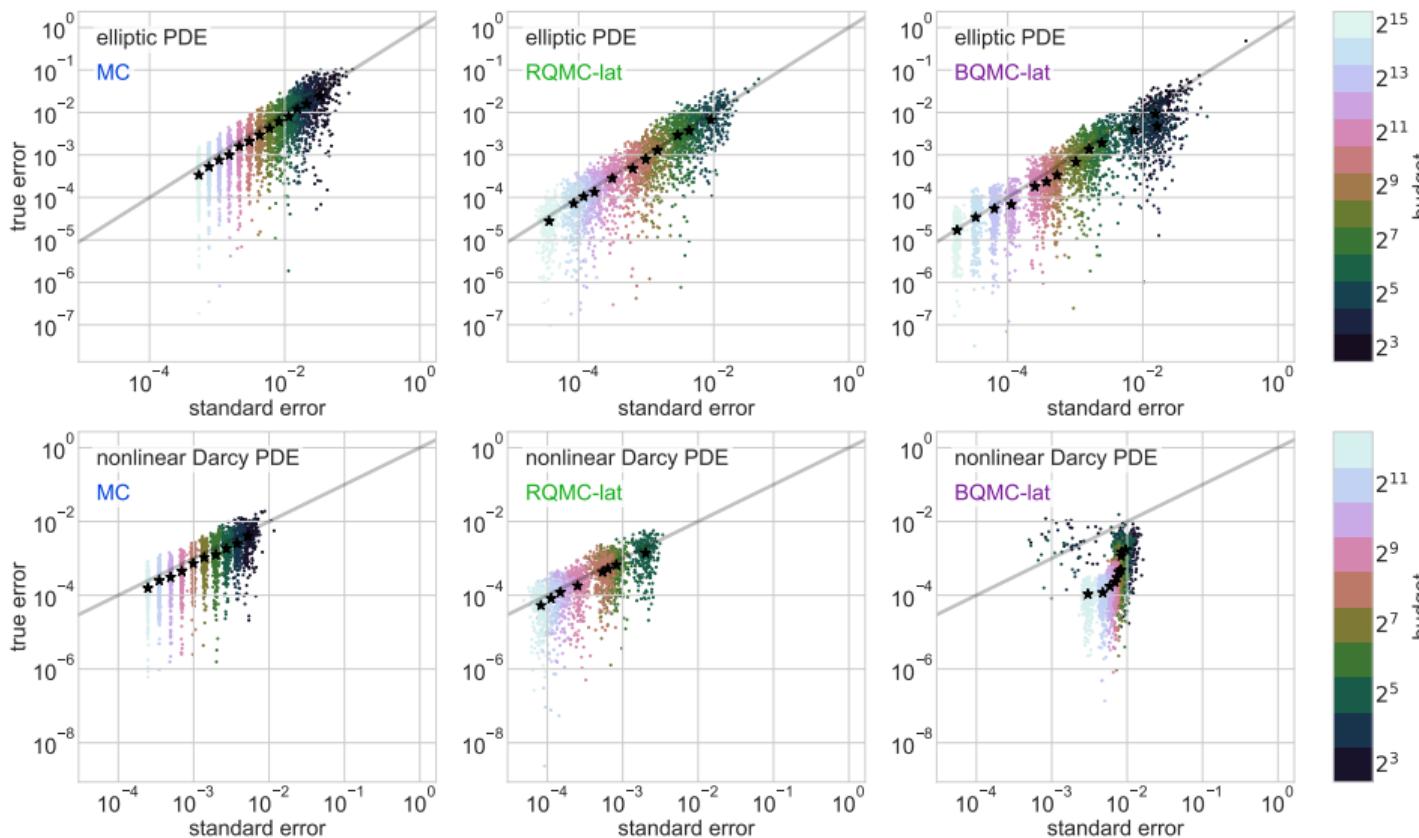
Single-Level Error Estimate Accuracy: LD Lattices



Multilevel Error Estimate Accuracy: LD Lattices

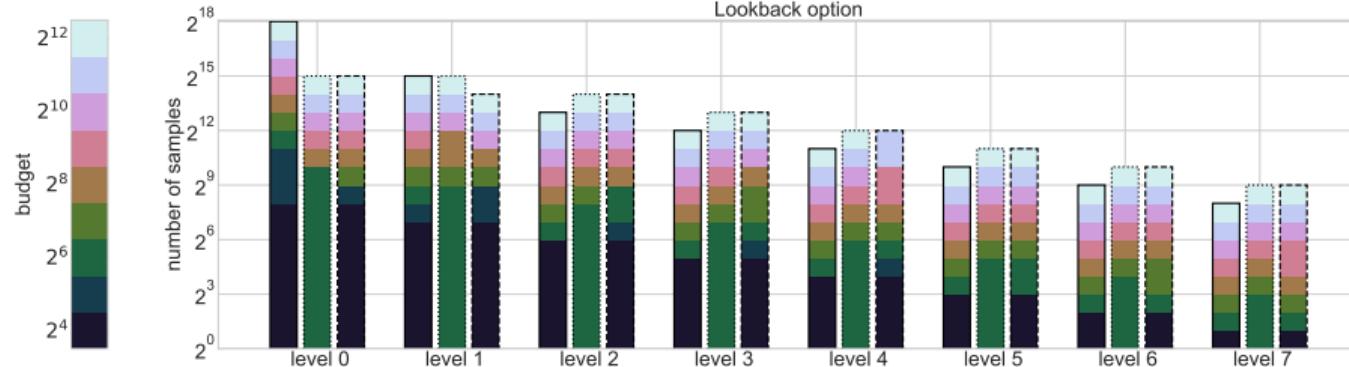
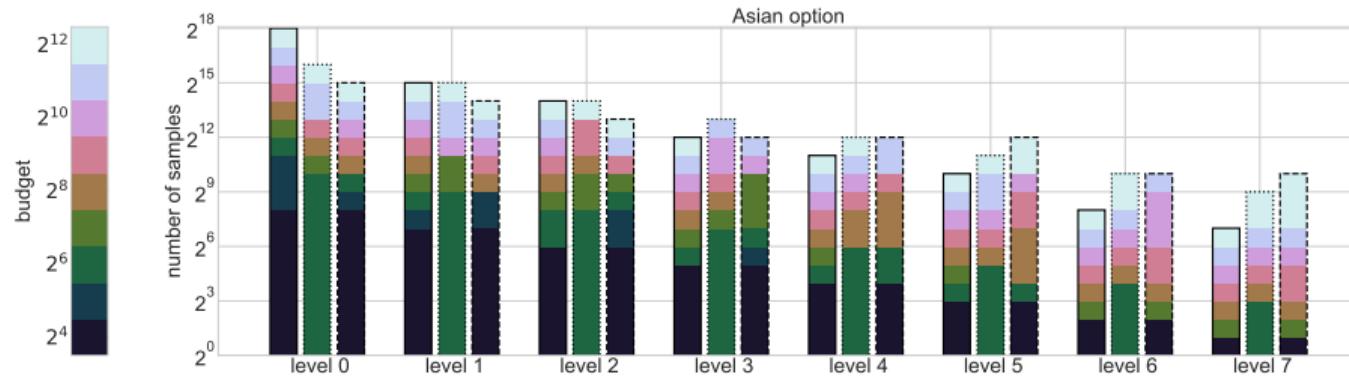


Multilevel Error Estimate Accuracy: LD Lattices



Multilevel Sample Allocation: LD Lattices

— MC ····· RQMC-lat - - - BQMC-lat



Multilevel Sample Allocation: LD Lattices

