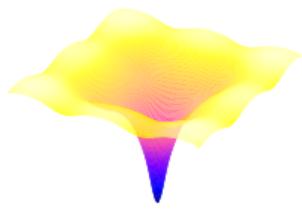


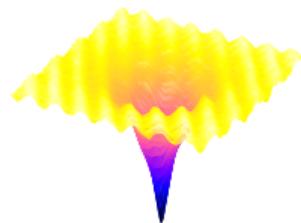
Fast Gaussian Process Regression

Aleksei G Sorokin^{1,2}, Pieterjan M Robbe², Fred J Hickernell¹
IIT, Department of Applied Math¹. Sandia National Lab².

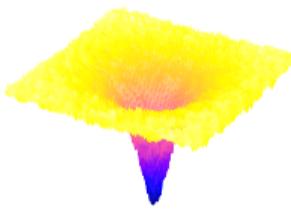
SE grid
 L_2 relative error = 3.71e-2
time = 2.24e0
average CI width = 2.45e-3
CI captured 0.1%



SI lattice
 L_2 relative error = 2.85e-2
time = 1.92e-3
average CI width = 3.71e0
CI captured 92.5%



DSI digital net
 L_2 relative error = 2.62e-2
time = 4.54e-3
average CI width = 2.74e0
CI captured 98.8%



FastGPs Software

`pip install fastgps` or visit alegresor.github.io/fastgps/

FastGPs Software

`pip install fastgps` or visit alegresor.github.io/fastgps/

A scalable Python software for fast Gaussian process regression (GPR) requiring only

- $\mathcal{O}(N \log N)$ computations (compared to typically $\mathcal{O}(N^3)$ cost), and
- $\mathcal{O}(N)$ memory (compare to classic $\mathcal{O}(N^2)$ memory requirements).

FastGPs Software

`pip install fastgps` or visit alegresor.github.io/fastgps/

A scalable Python software for fast Gaussian process regression (GPR) requiring only

- $\mathcal{O}(N \log N)$ computations (compared to typically $\mathcal{O}(N^3)$ cost), and
- $\mathcal{O}(N)$ memory (compare to classic $\mathcal{O}(N^2)$ memory requirements).

Fast GPR exploits Gram matrix structure present when pairing

- certain low discrepancy point sets (either lattices or digital nets) to
- certain shift invariant (SI) kernels (digitally SI (DSI) kernels for digital nets).

FastGPs Software

`pip install fastgps` or visit alegresor.github.io/fastgps/

A scalable Python software for fast Gaussian process regression (GPR) requiring only

- $\mathcal{O}(N \log N)$ computations (compared to typically $\mathcal{O}(N^3)$ cost), and
 - $\mathcal{O}(N)$ memory (compare to classic $\mathcal{O}(N^2)$ memory requirements).

Fast GPR exploits Gram matrix structure present when pairing

- certain low discrepancy point sets (either lattices or digital nets) to
 - certain shift invariant (SI) kernels (digitally SI (DSI) kernels for digital nets).

This requires

- control over the design of experiments (to use low discrepancy designs), and
 - using uncommon kernels (periodic SI kernels or discontinuous DS1 kernels).

FastGPs Software Features

`pip install fastgps` or visit alegresor.github.io/fastgps/

FastGPs Software Features

`pip install fastgps` or visit alegresor.github.io/fastgps/

1. **Kernel hyperparameter optimization** of marginal log likelihood (MLL), cross validation (CV), or generalized cross validation (GCV) loss.

FastGPs Software Features

`pip install fastgps` or visit alegresor.github.io/fastgps/

1. **Kernel hyperparameter optimization** of marginal log likelihood (MLL), cross validation (CV), or generalized cross validation (GCV) loss.
2. **Fast Bayesian cubature** for uncertainty quantification in Quasi-Monte Carlo.

FastGPs Software Features

`pip install fastgps` or visit alegresor.github.io/fastgps/

1. **Kernel hyperparameter optimization** of marginal log likelihood (MLL), cross validation (CV), or generalized cross validation (GCV) loss.
2. **Fast Bayesian cubature** for uncertainty quantification in Quasi-Monte Carlo.
3. **Fast multi-task GPR** with support for different sample sizes for each task.
Potentially useful for multi-fidelity simulations and Multilevel Monte Carlo.

FastGPs Software Features

`pip install fastgps` or visit alegresor.github.io/fastgps/

1. **Kernel hyperparameter optimization** of marginal log likelihood (MLL), cross validation (CV), or generalized cross validation (GCV) loss.
2. **Fast Bayesian cubature** for uncertainty quantification in Quasi-Monte Carlo.
3. **Fast multi-task GPR** with support for different sample sizes for each task.
Potentially useful for multi-fidelity simulations and Multilevel Monte Carlo.
4. **Batched GPR** for simultaneously modeling vector-output simulations.

FastGPs Software Features

`pip install fastgps` or visit alegresor.github.io/fastgps/

1. **Kernel hyperparameter optimization** of marginal log likelihood (MLL), cross validation (CV), or generalized cross validation (GCV) loss.
2. **Fast Bayesian cubature** for uncertainty quantification in Quasi-Monte Carlo.
3. **Fast multi-task GPR** with support for different sample sizes for each task.
Potentially useful for multi-fidelity simulations and Multilevel Monte Carlo.
4. **Batched GPR** for simultaneously modeling vector-output simulations.
5. **GPU support** enabled by the PyTorch stack .

FastGPs Software Features

`pip install fastgps` or visit alegresor.github.io/fastgps/

1. **Kernel hyperparameter optimization** of marginal log likelihood (MLL), cross validation (CV), or generalized cross validation (GCV) loss.
2. **Fast Bayesian cubature** for uncertainty quantification in Quasi-Monte Carlo.
3. **Fast multi-task GPR** with support for different sample sizes for each task.
Potentially useful for multi-fidelity simulations and Multilevel Monte Carlo.
4. **Batched GPR** for simultaneously modeling vector-output simulations.
5. **GPU support** enabled by the PyTorch stack .
6. **Flexible LD sequences and SI/DSI kernels** from the Quasi-Monte Carlo Python package QMCPy qmcssoftware.github.io/QMCSoftware/.

FastGPs Software Features

`pip install fastgps` or visit alegresor.github.io/fastgps/

1. **Kernel hyperparameter optimization** of marginal log likelihood (MLL), cross validation (CV), or generalized cross validation (GCV) loss.
2. **Fast Bayesian cubature** for uncertainty quantification in Quasi-Monte Carlo.
3. **Fast multi-task GPR** with support for different sample sizes for each task.
Potentially useful for multi-fidelity simulations and Multilevel Monte Carlo.
4. **Batched GPR** for simultaneously modeling vector-output simulations.
5. **GPU support** enabled by the PyTorch stack .
6. **Flexible LD sequences and SI/DSI kernels** from the Quasi-Monte Carlo Python package QMCPy qmcssoftware.github.io/QMCSSoftware/.
7. **Derivative-informed GPR** for simulations coupled with automatic differentiation.

FastGPs Software Features

`pip install fastgps` or visit alegresor.github.io/fastgps/

1. **Kernel hyperparameter optimization** of marginal log likelihood (MLL), cross validation (CV), or generalized cross validation (GCV) loss.
2. **Fast Bayesian cubature** for uncertainty quantification in Quasi-Monte Carlo.
3. **Fast multi-task GPR** with support for different sample sizes for each task.
Potentially useful for multi-fidelity simulations and Multilevel Monte Carlo.
4. **Batched GPR** for simultaneously modeling vector-output simulations.
5. **GPU support** enabled by the PyTorch stack .
6. **Flexible LD sequences and SI/DSI kernels** from the Quasi-Monte Carlo Python package QMCPy qmcssoftware.github.io/QMCSSoftware/.
7. **Derivative-informed GPR** for simulations coupled with automatic differentiation.
8. **Efficient variance projections** for non-greedy Bayesian optimization in MLMC.

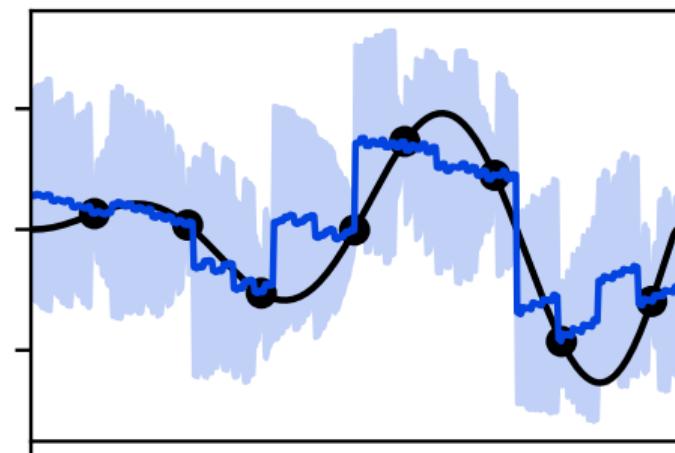
Gaussian Process Regression (GPR)

$f \sim \text{GP}(0, K)$ with posterior mean and covariance

$$\mathbb{E}[f(\mathbf{x})|\mathbf{X}, \mathbf{f}] = \mathbf{K}_{\mathbf{X}}(\mathbf{x}) \mathbf{K}^{-1} \mathbf{f}$$

$$\mathbb{V}[f(\mathbf{x})|\mathbf{X}, \mathbf{f}] = K(\mathbf{x}, \mathbf{x}) - \mathbf{K}_{\mathbf{X}}(\mathbf{x})^T \mathbf{K}^{-1} \mathbf{K}_{\mathbf{X}}(\mathbf{x})$$

- SPD $K : [0, 1]^d \times [0, 1]^d \rightarrow \mathbb{R}$
- Sampling locations $\mathbf{X} = \{\mathbf{x}_i\}_{i=0}^{N-1}$
- Sample values $\mathbf{f} = \{f(\mathbf{x}_i)\}_{i=0}^{N-1}$
- Kernel vector $\mathbf{K}_{\mathbf{X}}(\mathbf{x}) = \{K(\mathbf{x}, \mathbf{x}_i)\}_{i=0}^{N-1}$
- Gram matrix $\mathbf{K} = \{K(\mathbf{x}_i, \mathbf{x}_{i'})\}_{i,i'=0}^{N-1}$



Typically requires $\mathcal{O}(N^2)$ storage and $\mathcal{O}(N^3)$ computations

Fast GPs Pairing LD Points with Special Kernels

Recently explored in the context of fast Bayesian cubature [Rathinavel, 2019]

Fast GPs Pairing LD Points with Special Kernels

Recently explored in the context of fast Bayesian cubature [Rathinavel, 2019]

1. LD Lattices + Shift Invariant (SI) Kernels

- Give circulant Gram matrices $K = \{K(\mathbf{x}_i, \mathbf{x}_{i'})\}_{i,i'=0}^{N-1}$
- ∴ Eigendecomposition $K = V \Lambda \bar{V}$ where \bar{V} is the DFT matrix → FFT in $\mathcal{O}(N \log N)$
- [Rathinavel and Hickernell, 2019]

Fast GPs Pairing LD Points with Special Kernels

Recently explored in the context of fast Bayesian cubature [Rathinavel, 2019]

1. LD Lattices + Shift Invariant (SI) Kernels

- Give circulant Gram matrices $K = \{K(\mathbf{x}_i, \mathbf{x}_{i'})\}_{i,i'=0}^{N-1}$
- ∴ Eigendecomp $K = V \Lambda \bar{V}$ where \bar{V} is the DFT matrix → FFT in $\mathcal{O}(N \log N)$
- [Rathinavel and Hickernell, 2019]

2. LD Digital Nets + Digitally Shift Invariant (DSI) Kernels

- Give Recursive Symmetric Block Toeplitz (RSBT) Gram matrices K
- ∴ Eigendecomp $K = V \Lambda \bar{V}$ where \bar{V} is the Hadamard matrix → FWHT in $\mathcal{O}(N \log N)$
- [Rathinavel and Hickernell, 2022]

Fast GPs Pairing LD Points with Special Kernels

Recently explored in the context of fast Bayesian cubature [Rathinavel, 2019]

1. LD Lattices + Shift Invariant (SI) Kernels

- Give circulant Gram matrices $K = \{K(\mathbf{x}_i, \mathbf{x}_{i'})\}_{i,i'=0}^{N-1}$
- ∴ Eigendecomp $K = V \Lambda \bar{V}$ where \bar{V} is the DFT matrix → FFT in $\mathcal{O}(N \log N)$
- [Rathinavel and Hickernell, 2019]

2. LD Digital Nets + Digitally Shift Invariant (DSI) Kernels

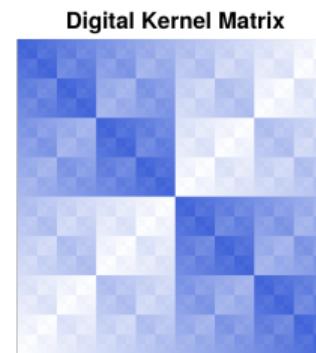
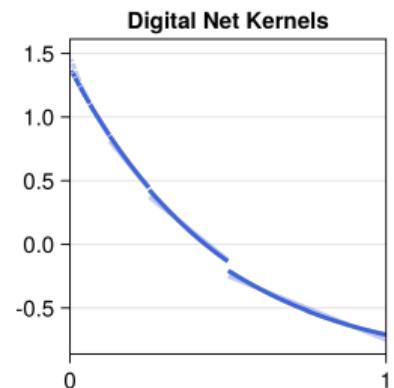
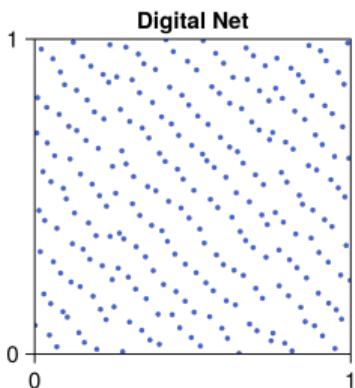
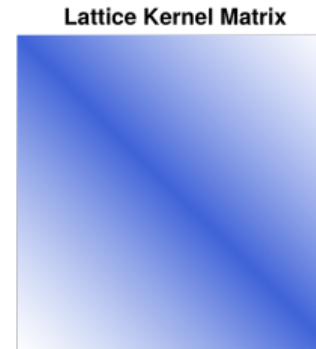
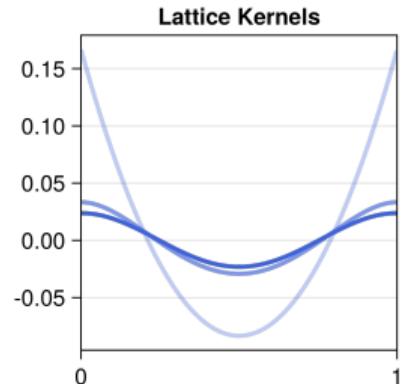
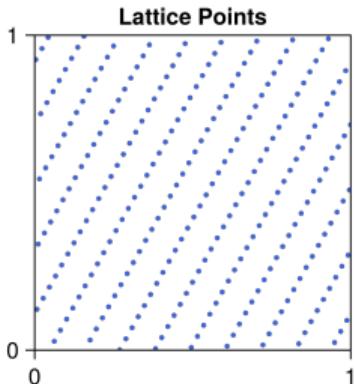
- Give Recursive Symmetric Block Toeplitz (RSBT) Gram matrices K
- ∴ Eigendecomp $K = V \Lambda \bar{V}$ where \bar{V} is the Hadamard matrix → FWHT in $\mathcal{O}(N \log N)$
- [Rathinavel and Hickernell, 2022]

Let $v_1 = \mathbf{1}/\sqrt{N}$ and k_1 be the first columns of \bar{V} and K respectively:

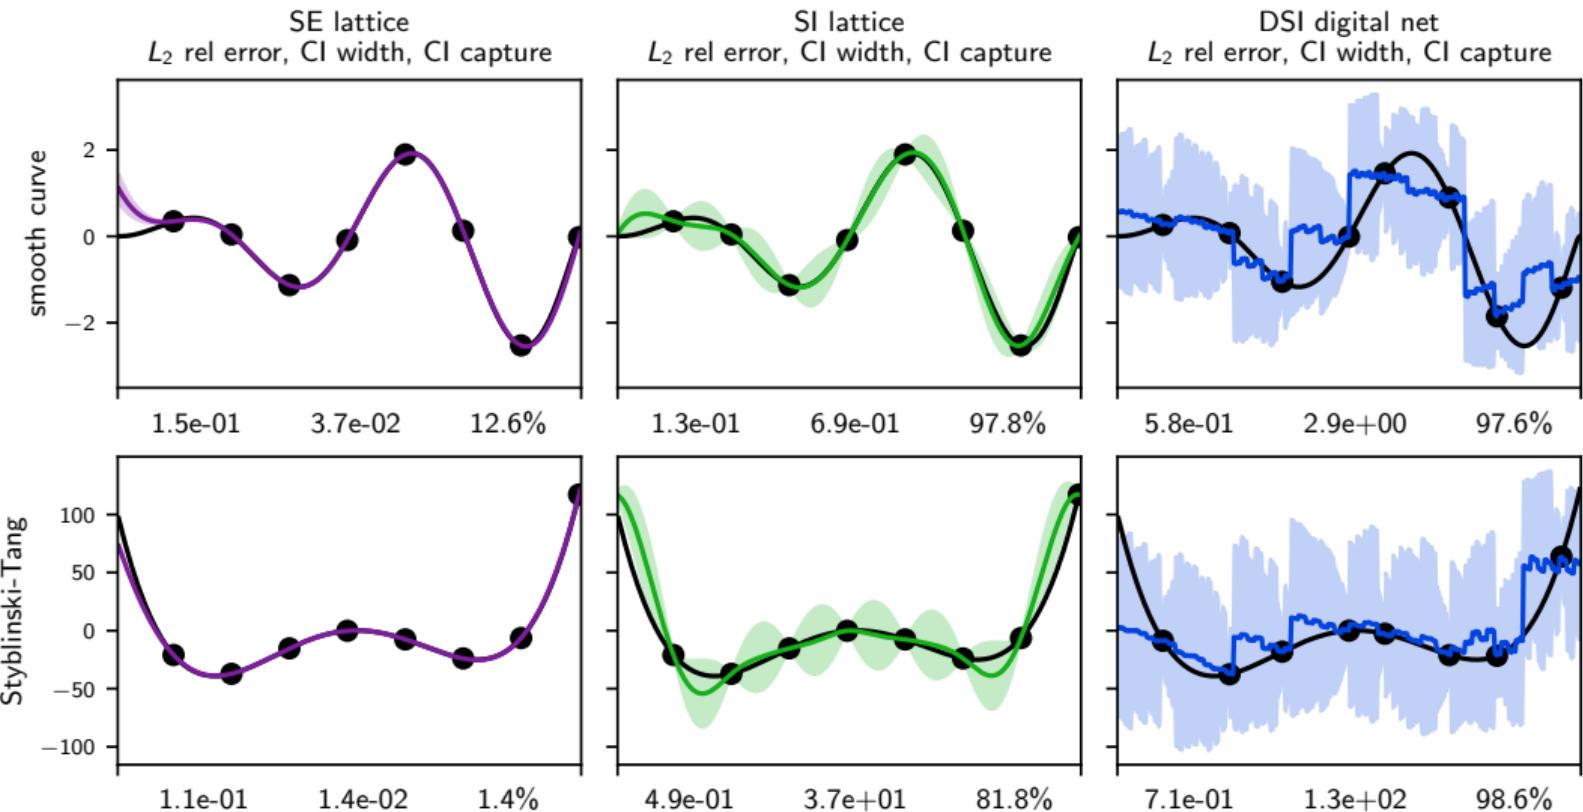
$$\lambda := \Lambda \mathbf{1} = \sqrt{N} \Lambda \bar{v}_1 = \sqrt{N} \bar{V} V \Lambda \bar{v}_1 = \sqrt{N} \bar{V} k_1$$

- Ka , $K^{-1}a$, and $|K|$ can all be computed in $\mathcal{O}(N \log N)$ computations
- Only requires evaluating and storing the first column of K

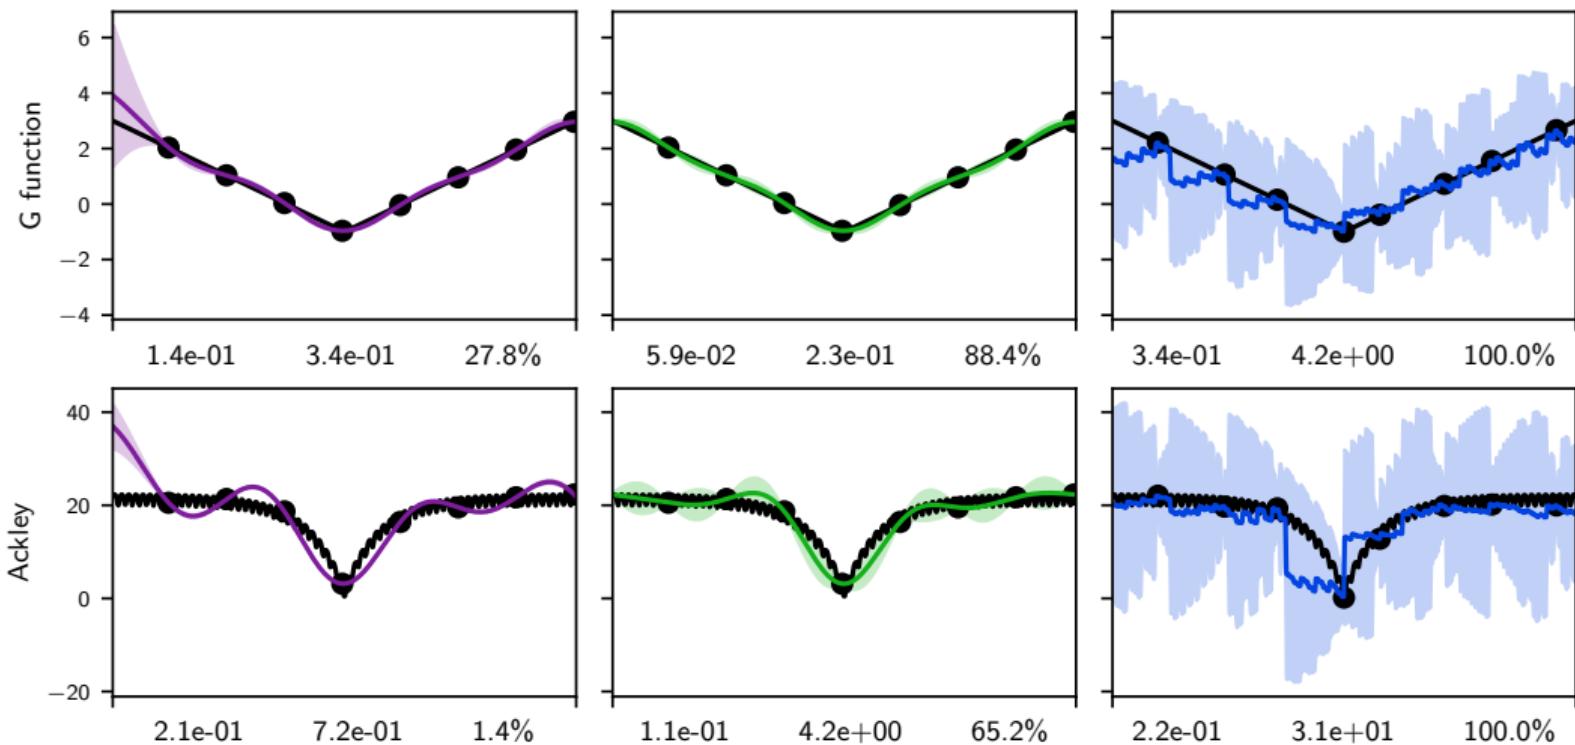
Lattices + SI $K =$ Circulant K and Digital Nets + DSI $K =$ RSBT K



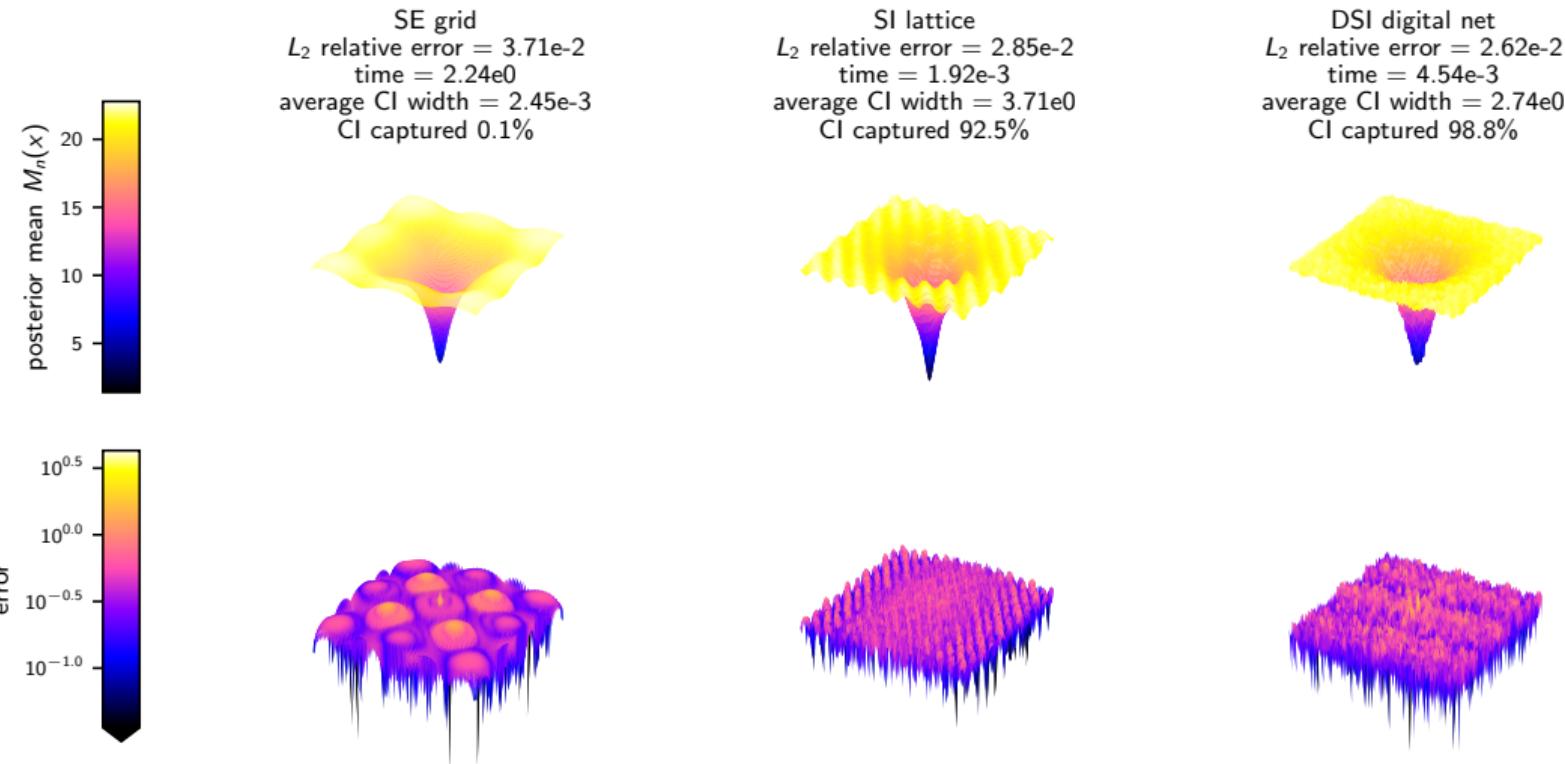
Examples in One Dimension [Surjanovic and Bingham], $N = 8$



Examples in One Dimension [Surjanovic and Bingham], $N = 8$



Ackley Function in Two Dimensions [Surjanovic and Bingham], $N = 4096$



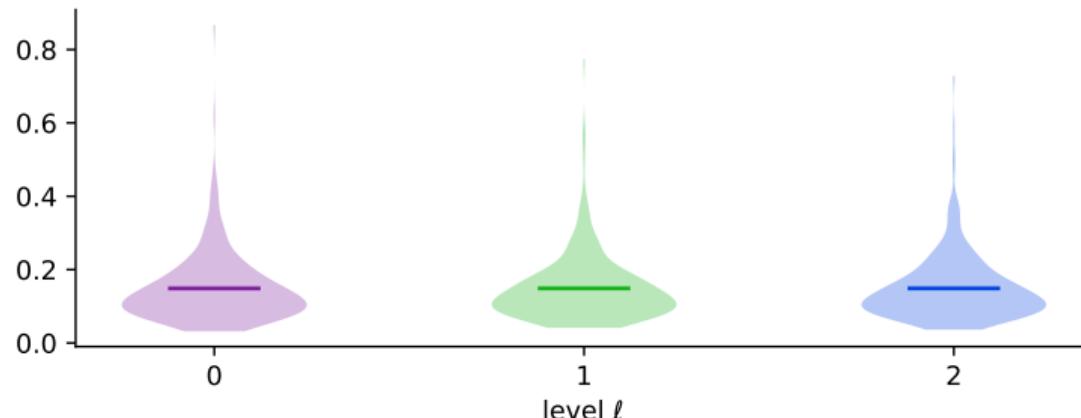
Multilevel (Multi-Task) Modeling

Given a multilevel simulation

$$f : \{1, \dots, L\} \times [0, 1]^d \rightarrow \mathbb{R},$$

we want to model $f(L, \cdot)$, the true (maximum-fidelity) simulation.

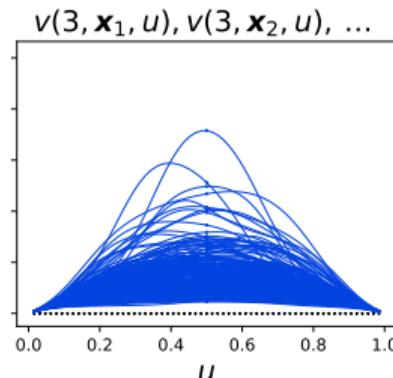
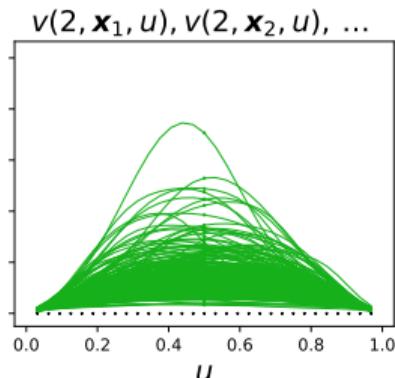
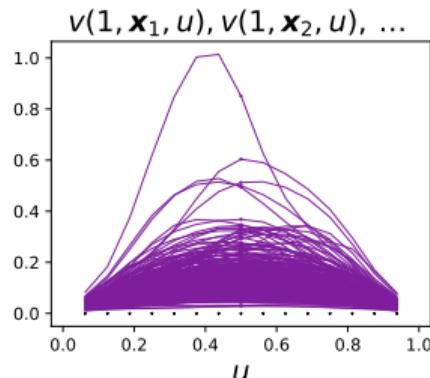
- $f(\ell, \mathbf{x})$ simulates at level $\ell \in \{1, \dots, L\}$ and with parameters $\mathbf{x} \in [0, 1]^d$
- Cost C_ℓ typically greater on higher levels
- $f(1, \cdot), f(2, \cdot), \dots, f(L, \cdot)$ are typically highly correlated



Numerical Solutions of PDEs with Random Coefficients

$$f(\ell, \mathbf{x}) = \mathcal{F}(v(\ell, \mathbf{x}, \cdot))$$

- $v : \{1, \dots, L\} \times [0, 1]^d \times \Omega$ is the numerical solution to the PDE
- \mathbf{x} represent random coefficients, e.g. coefficients in a Karhunen–Loève expansion
- ℓ controls the fidelity of the numerical solver e.g. the mesh width is $2^{-\ell}$
- \mathcal{F} is a (possibly non-linear) functional of the PDE solution, e.g.,
 - $\mathcal{F}(v(\ell, \mathbf{x}, \cdot)) = \mathbb{E}[v(\ell, \mathbf{x}, \mathbf{U})]$ where $\mathbf{U} \sim \mathcal{U}(\Omega)$, or
 - $\mathcal{F}(v(\ell, \mathbf{x}, \cdot)) = v(\ell, \mathbf{x}, u)$ for some $u \in \Omega$, e.g., $u = 1/2$ shown below



Multilevel (Quasi-)Monte Carlo without Replications

$$\mathbb{E}[f(L, \mathbf{X})] = \sum_{\ell=1}^L \mathbb{E}\underbrace{[f(\ell, \mathbf{X}) - f(\ell-1, \mathbf{X})]}_{\Delta_\ell(\mathbf{X})}, \quad \mathbf{X} \sim \mathcal{U}[0,1]^d, \quad f(0, \cdot) = 0$$

Multilevel (Quasi-)Monte Carlo without Replications

$$\mathbb{E}[f(L, \mathbf{X})] = \sum_{\ell=1}^L \mathbb{E}\underbrace{[f(\ell, \mathbf{X}) - f(\ell-1, \mathbf{X})]}_{\Delta_\ell(\mathbf{X})}, \quad \mathbf{X} \sim \mathcal{U}[0,1]^d, \quad f(0, \cdot) = 0$$

MLMC $\mathbb{E}[\Delta_\ell(\mathbf{X})] \approx \frac{1}{N_\ell} \sum_{i=0}^{N_\ell-1} \Delta_\ell(\mathbf{x}_i^\ell)$ for IID $\mathbf{x}_0^\ell, \dots, \mathbf{x}_{N_\ell-1}^\ell \sim \mathcal{U}[0,1]^d$

- $N_\ell \propto \sqrt{\mathbb{V}[\Delta_\ell(\mathbf{X})]/C_\ell}$ chosen to optimally minimize error [Giles, 2008]

Multilevel (Quasi-)Monte Carlo without Replications

$$\mathbb{E}[f(L, \mathbf{X})] = \sum_{\ell=1}^L \mathbb{E}\underbrace{[f(\ell, \mathbf{X}) - f(\ell-1, \mathbf{X})]}_{\Delta_\ell(\mathbf{X})}, \quad \mathbf{X} \sim \mathcal{U}[0,1]^d, \quad f(0, \cdot) = 0$$

MLMC $\mathbb{E}[\Delta_\ell(\mathbf{X})] \approx \frac{1}{N_\ell} \sum_{i=0}^{N_\ell-1} \Delta_\ell(\mathbf{x}_i^\ell)$ for IID $\mathbf{x}_0^\ell, \dots, \mathbf{x}_{N_\ell-1}^\ell \sim \mathcal{U}[0,1]^d$

- $N_\ell \propto \sqrt{\mathbb{V}[\Delta_\ell(\mathbf{X})]/C_\ell}$ chosen to optimally minimize error [Giles, 2008]

R-MLQMC $\mathbb{E}[\Delta_\ell(\mathbf{X})] \approx \frac{1}{R} \sum_{r=1}^R \frac{1}{N_\ell} \sum_{i=0}^{N_\ell-1} \Delta_\ell(\mathbf{x}_{ri}^\ell)$ [Giles and Waterhouse, 2009]

- IID randomizations of low discrepancy pointsets $\{\mathbf{x}_{1i}^\ell\}_{i=0}^{N_\ell-1}, \dots, \{\mathbf{x}_{Ri}^\ell\}_{i=0}^{N_\ell-1}$.
- N_ℓ greedily doubled on level where $\mathbb{V}[\Delta_\ell(\mathbf{X})]/(N_\ell C_\ell)$ the largest
- Variance approximated by sample variance across R sub-means

Multilevel (Quasi-)Monte Carlo without Replications

$$\mathbb{E}[f(L, \mathbf{X})] = \sum_{\ell=1}^L \mathbb{E}\underbrace{[f(\ell, \mathbf{X}) - f(\ell-1, \mathbf{X})]}_{\Delta_\ell(\mathbf{X})}, \quad \mathbf{X} \sim \mathcal{U}[0,1]^d, \quad f(0, \cdot) = 0$$

MLMC $\mathbb{E}[\Delta_\ell(\mathbf{X})] \approx \frac{1}{N_\ell} \sum_{i=0}^{N_\ell-1} \Delta_\ell(\mathbf{x}_i^\ell)$ for IID $\mathbf{x}_0^\ell, \dots, \mathbf{x}_{N_\ell-1}^\ell \sim \mathcal{U}[0,1]^d$

- $N_\ell \propto \sqrt{\mathbb{V}[\Delta_\ell(\mathbf{X})]/C_\ell}$ chosen to optimally minimize error [Giles, 2008]

R-MLQMC $\mathbb{E}[\Delta_\ell(\mathbf{X})] \approx \frac{1}{R} \sum_{r=1}^R \frac{1}{N_\ell} \sum_{i=0}^{N_\ell-1} \Delta_\ell(\mathbf{x}_{ri}^\ell)$ [Giles and Waterhouse, 2009]

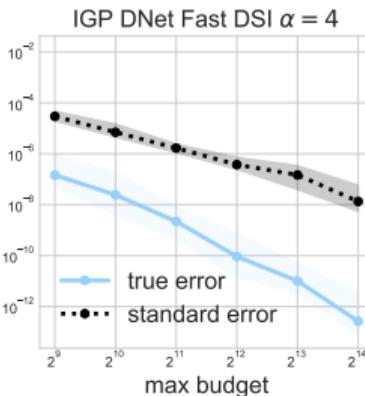
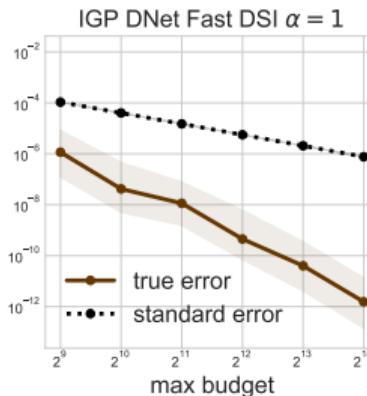
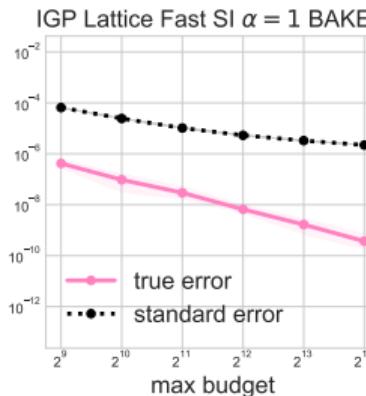
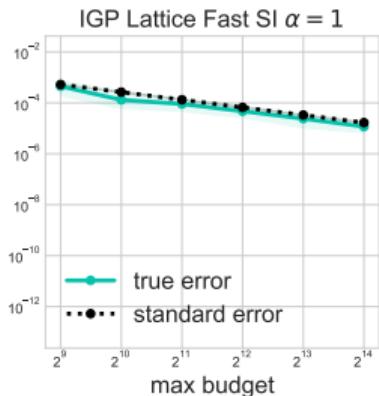
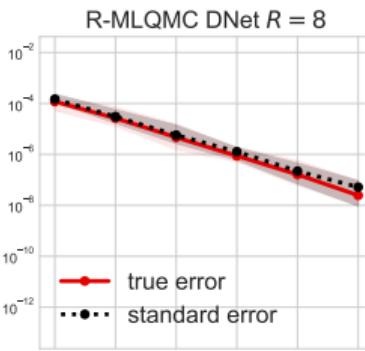
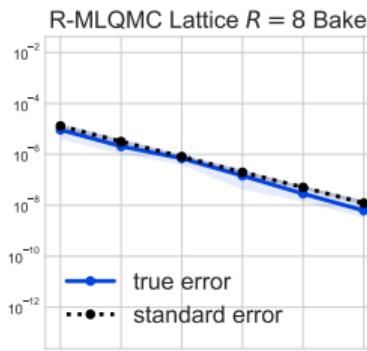
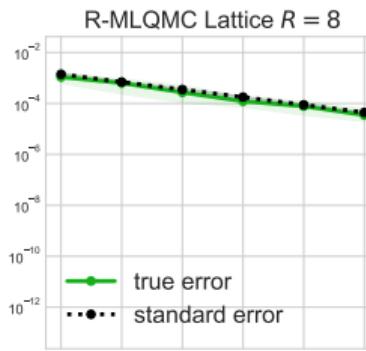
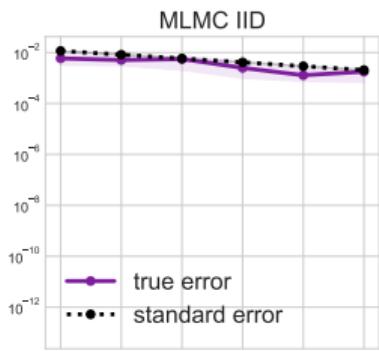
- IID randomizations of low discrepancy pointsets $\{\mathbf{x}_{1i}^\ell\}_{i=0}^{N_\ell-1}, \dots, \{\mathbf{x}_{Ri}^\ell\}_{i=0}^{N_\ell-1}$.
- N_ℓ greedily doubled on level where $\mathbb{V}[\Delta_\ell(\mathbf{X})]/(N_\ell C_\ell)$ the largest
- Variance approximated by sample variance across R sub-means

(IGP) Bayesian MLQMC without replications $\mathbb{E}[\Delta_\ell(\mathbf{X})] \approx \frac{1}{N_\ell} \sum_{i=0}^{N_\ell-1} \Delta_\ell(\mathbf{x}_i^\ell)$

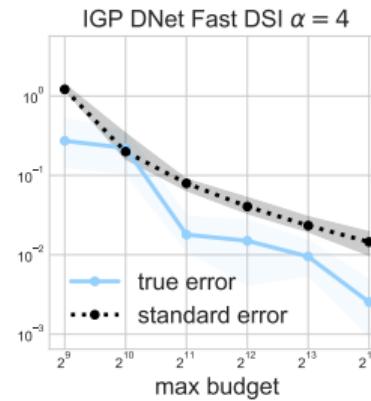
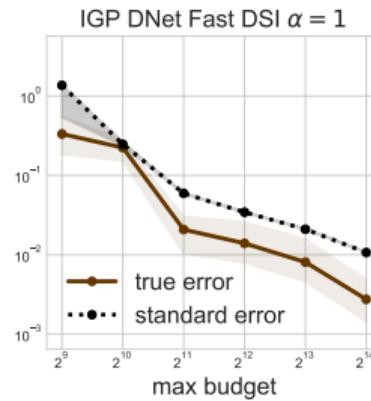
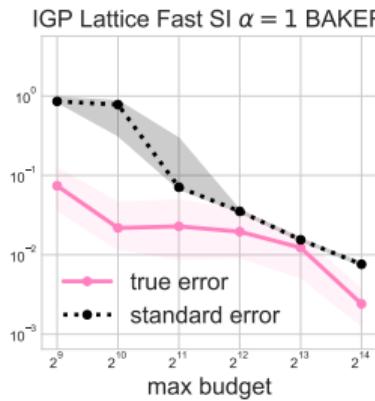
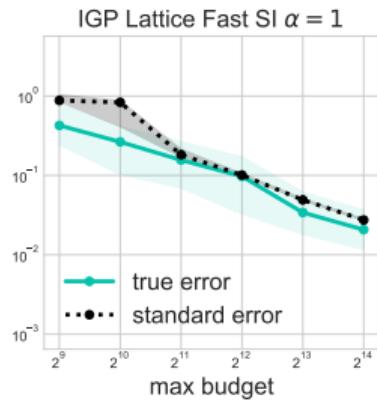
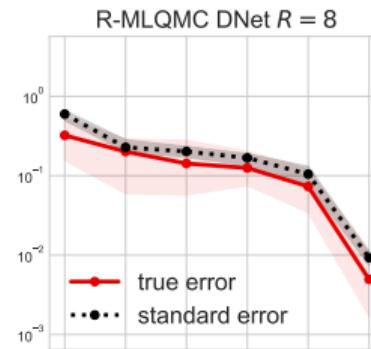
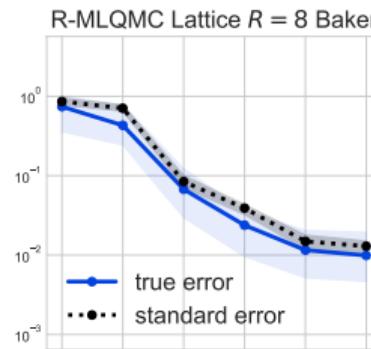
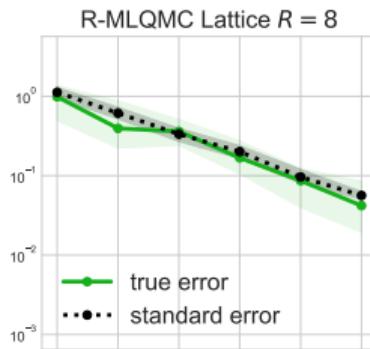
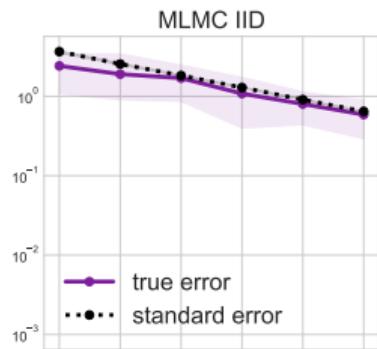
- $\{\mathbf{x}_i^\ell\}_{i=0}^{N_\ell-1}$ a single randomized low discrepancy pointset
- N_ℓ greedily doubled on level where $\mathbb{V}[\Delta_\ell(\mathbf{X})]/(N_\ell C_\ell)$ the largest
- Variance taken to be posterior variance with $\Delta_\ell \sim \text{GP}(0, K_\ell)$ (independent GPs)

$$f(\ell, (x_1, x_2)) = \sin(x_1) + 2^{-\ell} \cos(x_2)$$

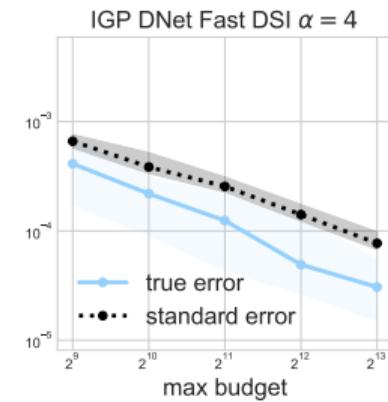
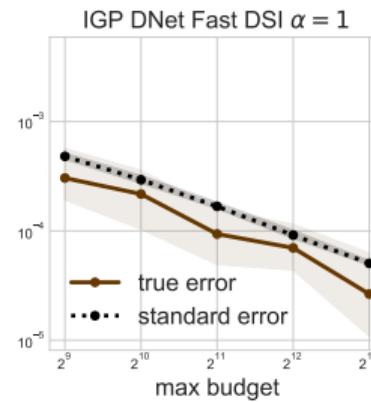
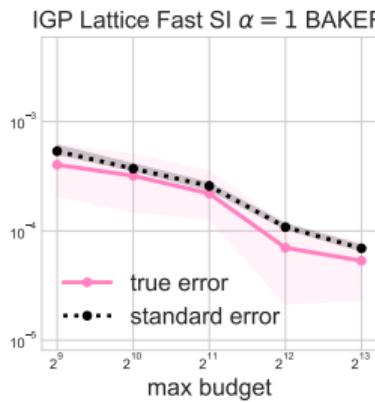
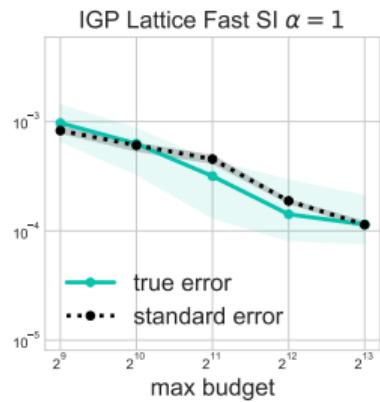
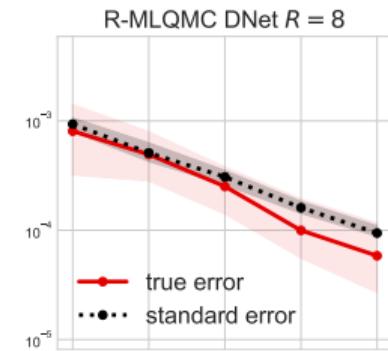
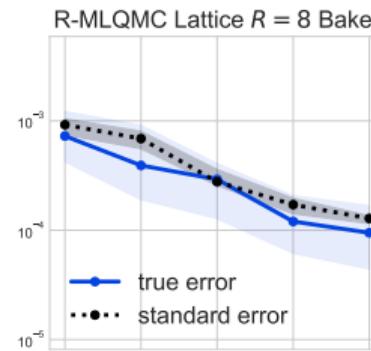
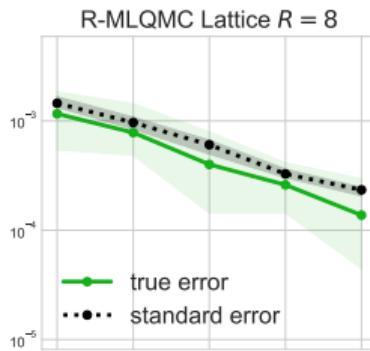
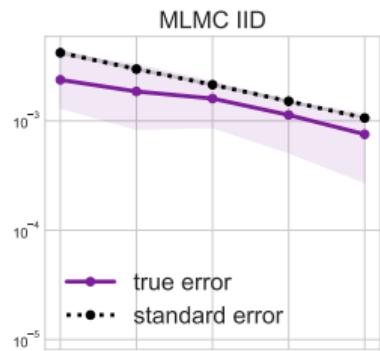
analytic problem with $d = 2$ dimensions and $L = 4$ levels with $T = 50$ trials



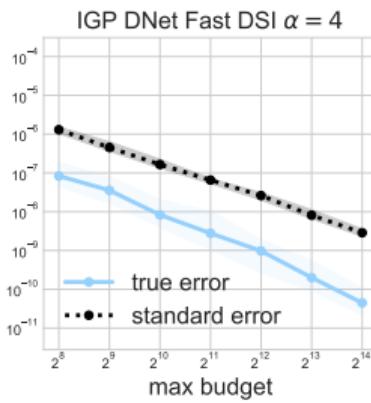
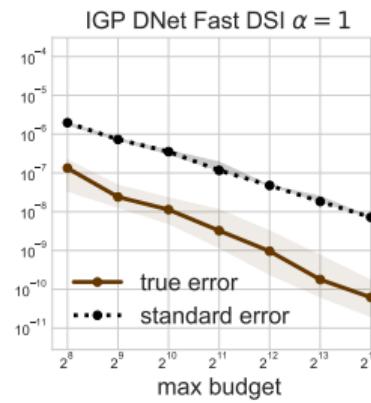
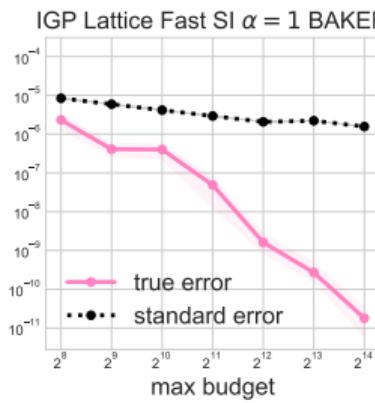
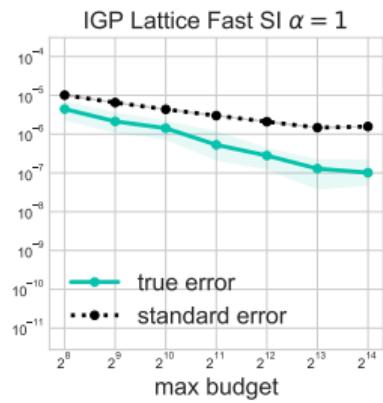
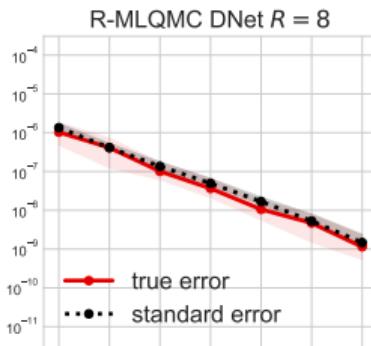
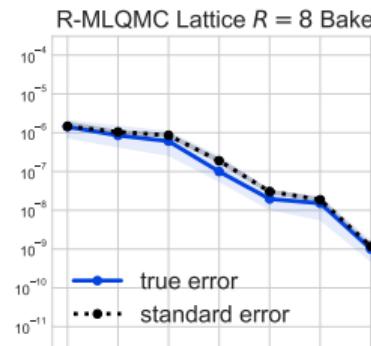
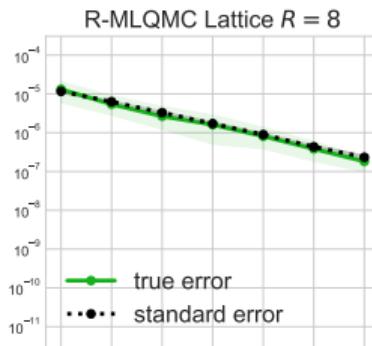
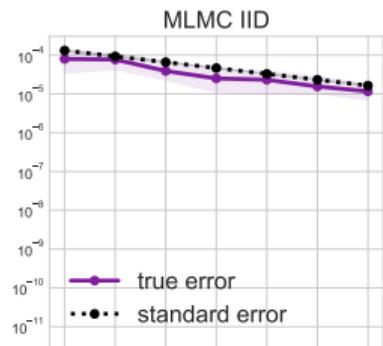
borehole problem with $d = 8$ dimensions and $L = 2$ levels with $T = 50$ trials

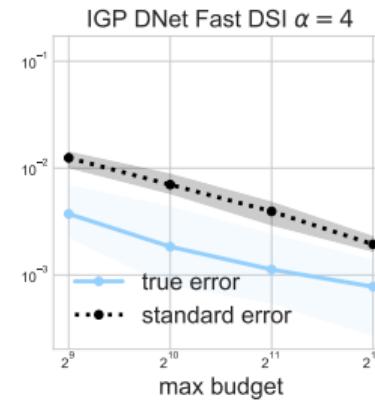
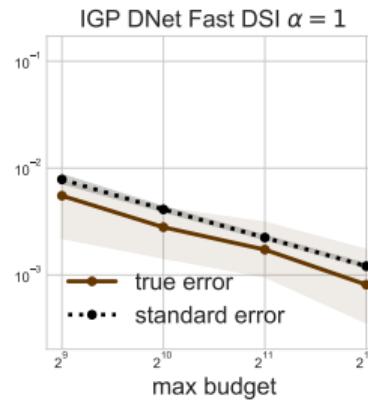
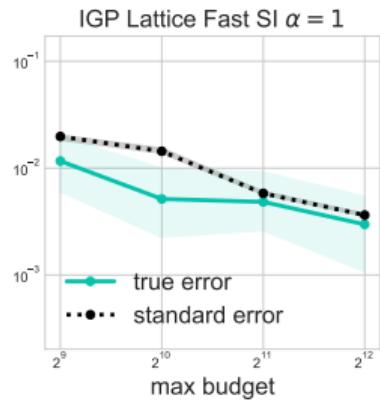
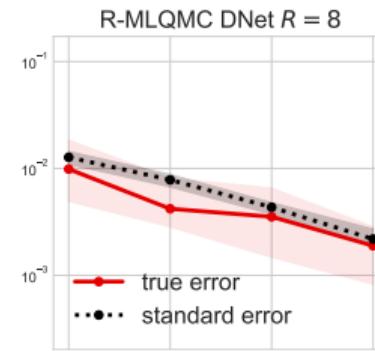
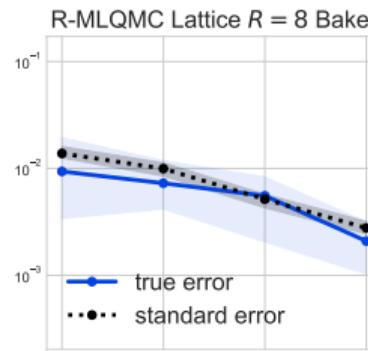
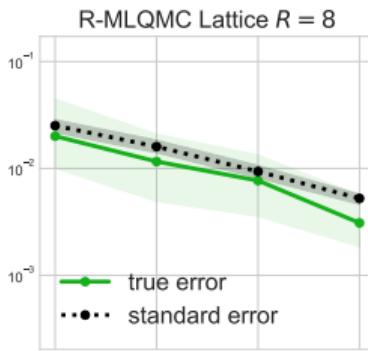
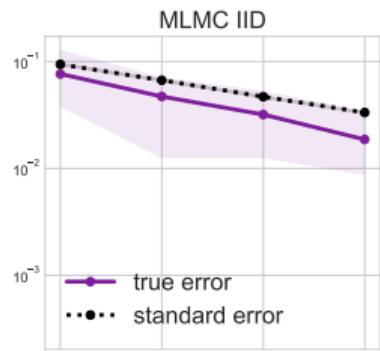


elliptic problem with $d = 8$ dimensions and $L = 4$ levels with $T = 50$ trials



steady state diffusion 1d problem with $d = 9$ dimensions and $L = 5$ levels with $T = 50$ trials



Asian option problem with $d = 16$ dimensions and $L = 8$ levels with $T = 50$ trials

Multi-Task Gaussian Processes

$$f(\ell, \mathbf{x}) \sim \text{GP}(0, K)$$

$N = N_1 + \dots + N_L$ sampling locations $\mathcal{D} = \mathcal{D}_1 \cup \dots \cup \mathcal{D}_L$ where $\mathcal{D}_\ell = \{(\ell, \mathbf{x}_i^\ell)\}_{i=1}^{N_\ell}$.
Posterior mean and covariance

$$\mathbb{E}[f(\ell, \mathbf{x})] = \mathbf{K}^T(\ell, \mathbf{x}) \mathbf{K}^{-1} \mathbf{f}$$

$$\mathbb{V}[f(\ell, \mathbf{x})] = K((\ell, \mathbf{x}), (\ell, \mathbf{x})) - \mathbf{K}^T(\ell, \mathbf{x}) \mathbf{K}^{-1} \mathbf{K}(\ell, \mathbf{x})$$

- $\mathbf{K}(\ell, \mathbf{x}) = K(\mathcal{D}, (\ell, \mathbf{x}))$ and $\mathbf{f} = f(\mathcal{D})$ are length N vectors
- $\mathbf{K} = K(\mathcal{D}, \mathcal{D}^T)$ is the $N \times N$ Gram matrix

Kernel K depends on hyperparameters θ e.g. global scale, lengthscale, etc.

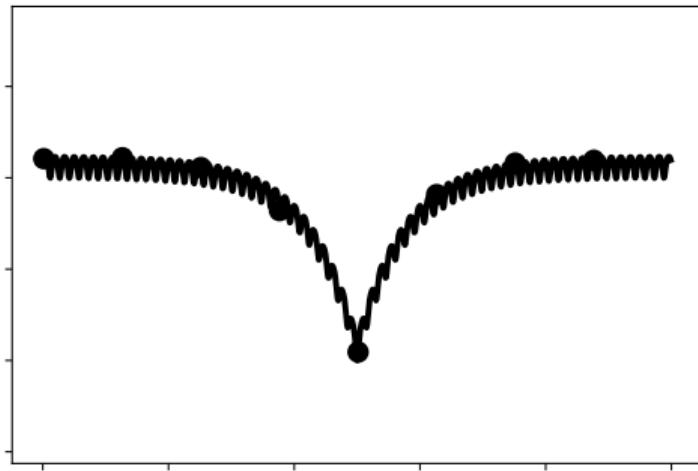
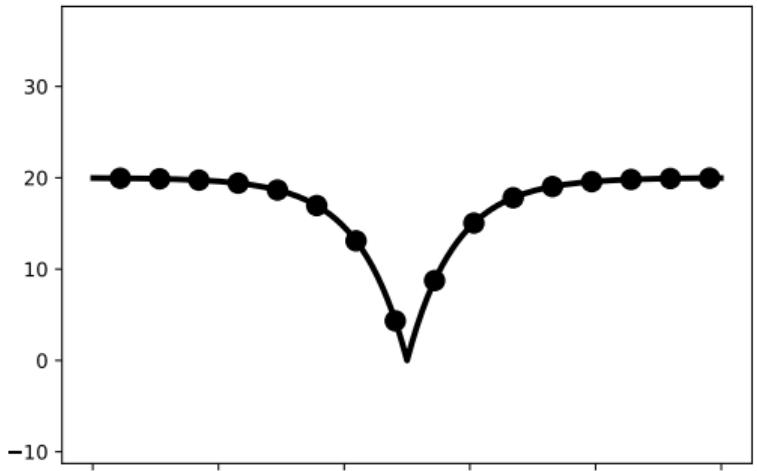
Hyperparameters θ often chosen to minimize negative marginal log likelihood (NMLL)

$$\text{NMLL} \propto \mathbf{f}^T \mathbf{K}^{-1} \mathbf{f} + \log|\mathbf{K}| + \log(2\pi)N$$

∴ Multi-Task GPs require computing $\mathbf{K}^{-1} \mathbf{f}$ and $\log|\mathbf{K}| \implies$ standard cost $\mathcal{O}(N^3)$

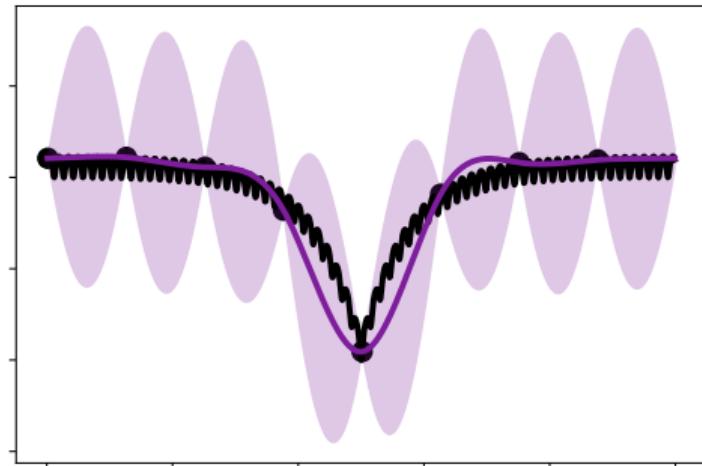
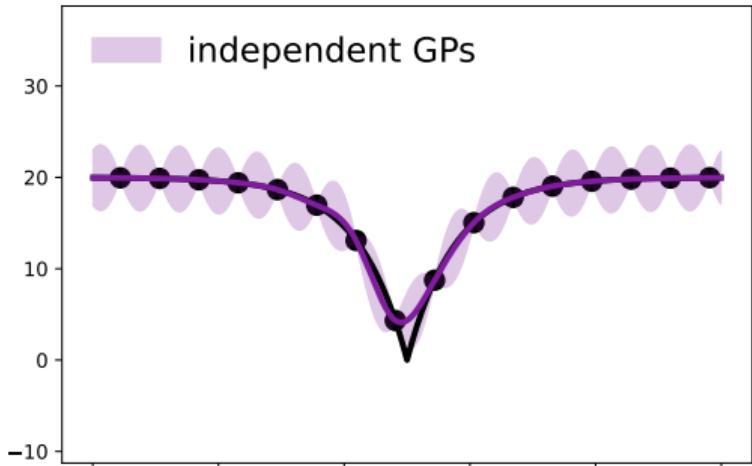
Independent GPs vs a Multi-Task GP Example

Low Fidelity Left, High Fidelity Right



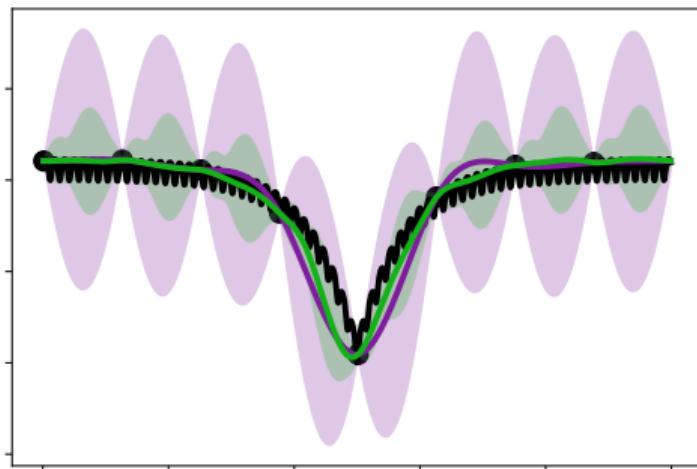
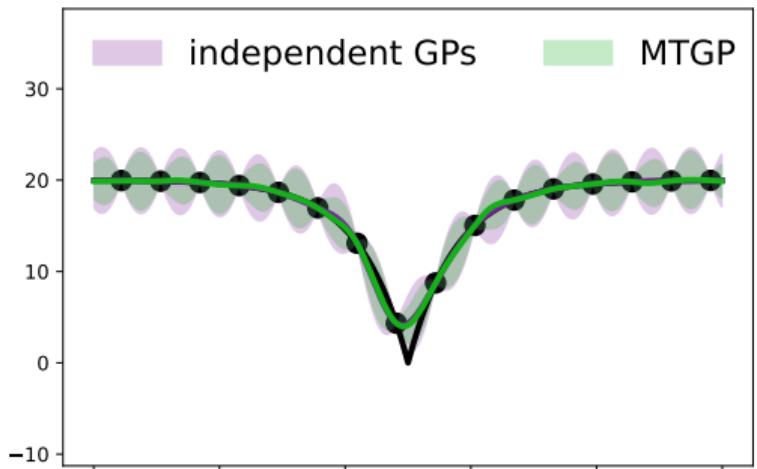
Independent GPs vs a Multi-Task GP Example

Low Fidelity Left, High Fidelity Right



Independent GPs vs a Multi-Task GP Example

Low Fidelity Left, High Fidelity Right



Product Kernels for Multi-Task GPs

Common to assume

$$K((\ell, \mathbf{x}), (\ell', \mathbf{x}')) = R(\ell, \ell') Q(\mathbf{x}, \mathbf{x}')$$

- $R : \{1, \dots, L\} \times \{1, \dots, L\} \rightarrow \mathbb{R}$ an SPD kernel over levels e.g.

$$\mathbf{R} = \{R(\ell, \ell')\}_{\ell, \ell'=1}^L = \mathbf{B}\mathbf{B}^T + \text{diag}(\boldsymbol{\nu}), \quad \boldsymbol{\nu} \in \mathbb{R}_+^L, \quad \mathbf{B} \in \mathbb{R}^{L \times r}, \quad \text{rank } r \leq L$$

- $Q : [0, 1]^d \times [0, 1]^d \rightarrow \mathbb{R}$ an SPD kernel over parameters

$$\mathbf{K} = \begin{pmatrix} \mathbf{K}_{11} & \cdots & \mathbf{K}_{1L} \\ \vdots & \ddots & \vdots \\ \overline{\mathbf{K}_{1L}} & \cdots & \mathbf{K}_{LL} \end{pmatrix} = \begin{pmatrix} R_{11}\mathbf{Q}_{11} & \cdots & R_{1L}\mathbf{Q}_{1L} \\ \vdots & \ddots & \vdots \\ \overline{R_{1L}\mathbf{Q}_{1L}} & \cdots & R_{LL}\mathbf{Q}_{LL} \end{pmatrix}$$

- $R_{\ell\ell'} = R(\ell, \ell')$ is a scalar
- $\mathbf{Q}_{\ell\ell'} = Q(\mathcal{D}_\ell, \mathcal{D}_{\ell'}^T)$ is an $N_\ell \times N_{\ell'}$ Gram matrix

Fast Multi-Task GPs

$$\mathbf{K} = \begin{pmatrix} R_{11}\mathbf{Q}_{11} & \cdots & R_{1L}\mathbf{Q}_{1L} \\ \vdots & \ddots & \vdots \\ R_{1L}\mathbf{Q}_{1L} & \cdots & R_{LL}\mathbf{Q}_{LL} \end{pmatrix}$$

Idea: Force “nice” structure in $\mathbf{Q}_{\ell\ell'}$ through special pairings of $\mathbf{X}_\ell = \{\mathbf{x}_i^\ell\}_{i=1}^{N_\ell}$ and \mathbf{Q}

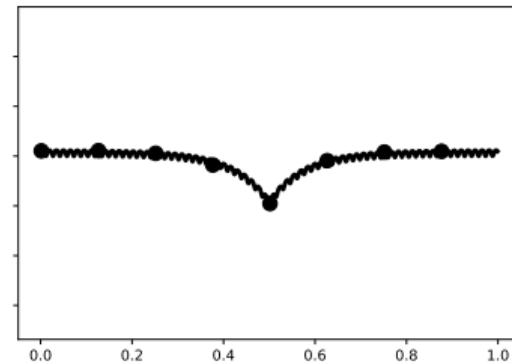
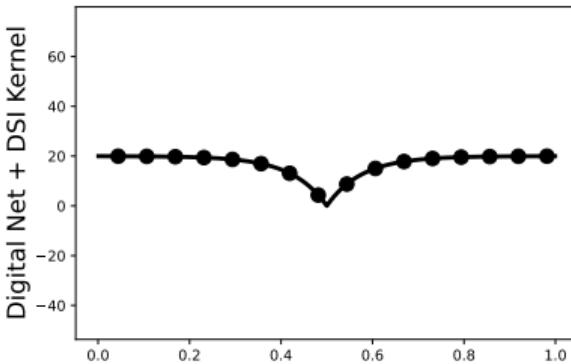
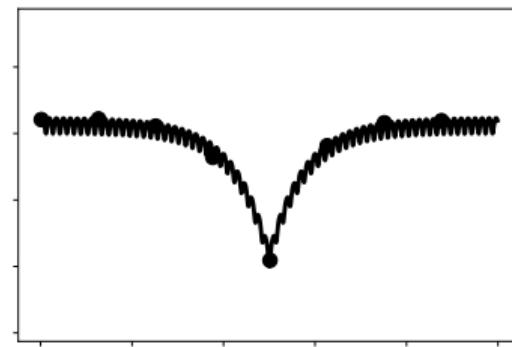
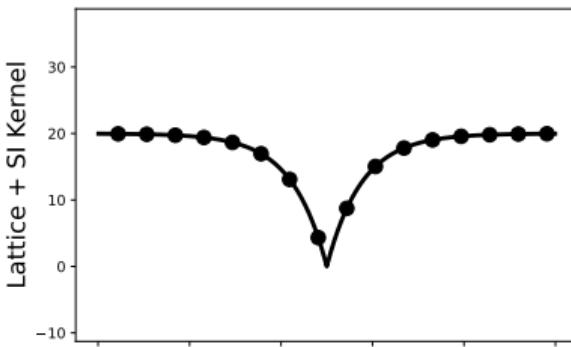
1. \mathbf{X}_ℓ a lattice and \mathbf{Q} a shift invariant (SI) kernel $\implies \mathbf{Q}_{\ell\ell'}$ block circulant
2. \mathbf{X}_ℓ a (base 2) digital net and \mathbf{Q} a digitally SI (DSI) kernel $\implies \mathbf{Q}_{\ell\ell'}$ block RSBT

Technicalities

- Lattices and digital nets require sample sizes $N_\ell = 2^{m_\ell}$
- Lattices $\mathbf{X}_1, \dots, \mathbf{X}_L$: same generating vector, possibly different random shifts
- Circulant matrices diagonalizable by FFT
- Digital nets $\mathbf{X}_1, \dots, \mathbf{X}_L$: same generating matrices, possibly different digital shifts
- RSBT matrices diagonalizable by FWHT

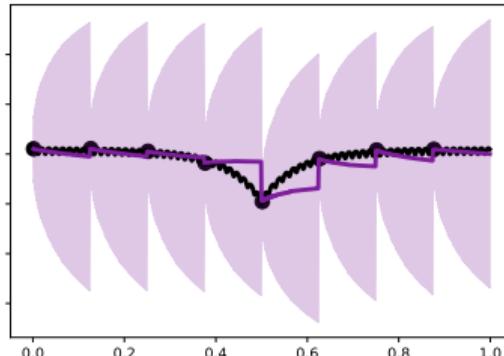
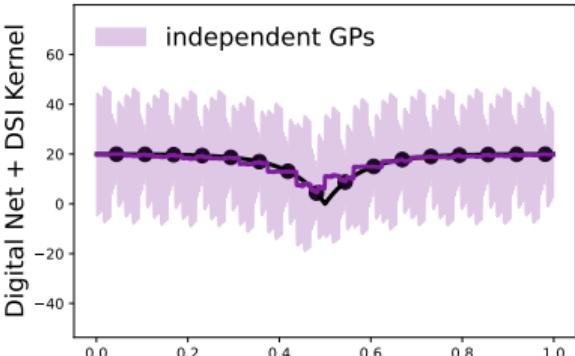
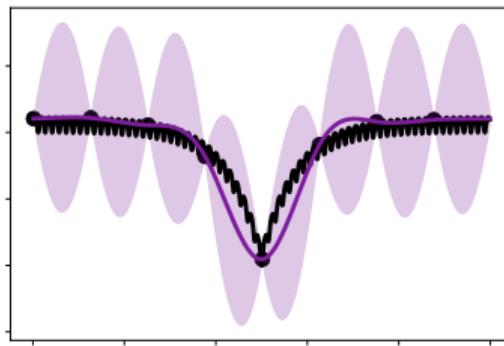
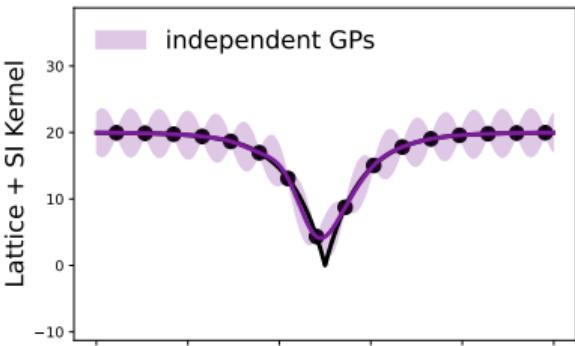
Independent Fast GPs vs Fast Multi-Task GPs Example

Low Fidelity Left, High Fidelity Right



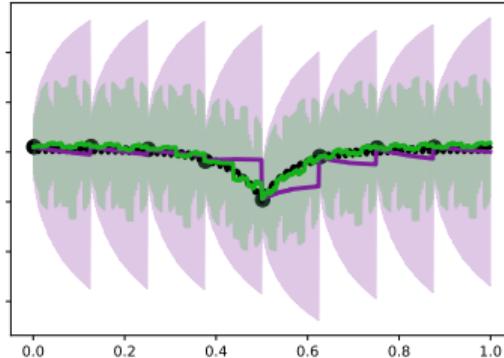
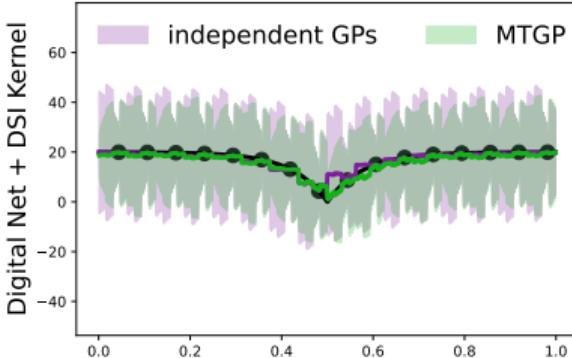
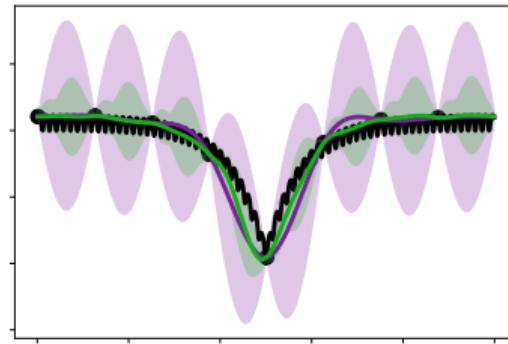
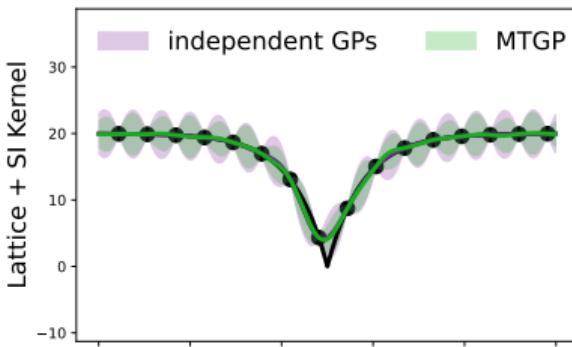
Independent Fast GPs vs Fast Multi-Task GPs Example

Low Fidelity Left, High Fidelity Right



Independent Fast GPs vs Fast Multi-Task GPs Example

Low Fidelity Left, High Fidelity Right



Fast Multi-Task GPs Continued

$$\mathbf{K}_{\ell\ell'} = \mathbf{R}_{\ell\ell'} \mathbf{Q}_{\ell\ell'} = \mathbf{V}_{m_\ell} \boldsymbol{\Sigma}_{\ell\ell'} \overline{\mathbf{V}_{m_{\ell'}}}$$

- $\overline{\mathbf{V}_m}$ a $2^m \times 2^m$ *fast transform matrix*
 1. Lattice X_ℓ with SI Q makes $\overline{\mathbf{V}_{m_\ell}}$ the Fast Fourier Transform
 2. Digital Net X_ℓ with DSI Q makes $\overline{\mathbf{V}_{m_\ell}}$ the Fast Walsh Hadamard Transform
- $\mathbf{V}_m \mathbf{a}$ and $\overline{\mathbf{V}_m} \mathbf{a}$ both cost only $\mathcal{O}(m2^m)$ to compute
- The first column of $\overline{\mathbf{V}_m}$ is $\mathbf{1}_m / \sqrt{2^m}$
- $\boldsymbol{\Sigma}_{\ell\ell'}$ a diagonal block matrix characterized by

$$\sigma_{\ell\ell'} = \boldsymbol{\Sigma}_{\ell\ell'} \mathbf{1}_{m_{\ell'}} = \sqrt{2^{m_{\ell'}}} \overline{\mathbf{V}_{m_\ell}} \mathbf{k}_{\ell\ell',1}$$

where $\mathbf{k}_{\ell\ell',1}$ is the first column of $\mathbf{K}_{\ell\ell'}$ and we assume $m_\ell \geq m_{\ell'}$

Fast Multi-Task GPs NMLL

$$\mathbf{K} = \begin{pmatrix} \mathbf{V}_{m_1} & & \\ & \ddots & \\ & & \mathbf{V}_{m_L} \end{pmatrix} \begin{pmatrix} \Sigma_{11} & \cdots & \Sigma_{1L} \\ \vdots & \ddots & \vdots \\ \overline{\Sigma_{1L}} & \cdots & \Sigma_{LL} \end{pmatrix} \begin{pmatrix} \sqrt{\mathbf{V}_{m_1}} & & \\ & \ddots & \\ & & \sqrt{\mathbf{V}_{m_L}} \end{pmatrix} =: \mathbf{V}\Sigma\mathbf{V}$$

$$\text{NMLL} \propto \hat{\mathbf{f}}^T \Sigma^{-1} \hat{\mathbf{f}} + \log|\Sigma| + \log(2\pi)N, \quad \hat{\mathbf{f}} = \bar{\mathbf{V}}\mathbf{f}$$

∴ Fast Multi-Task GP fitting requires computing $\Sigma^{-1}\hat{\mathbf{f}}$ and $\log|\Sigma|$

For example, if $N_1 = 8$, $N_2 = 4$, and $N_3 = 2$ then Σ has the following structure

$$\Sigma = \left[\begin{array}{c|c|c|c} \cdot & & & \cdot \\ \hline & \cdot & & \cdot \\ \hline & & \cdot & \cdot \\ \hline & & & \cdot \\ \hline \cdot & & & \cdot \\ \hline & \cdot & & \cdot \\ \hline & & \cdot & \cdot \\ \hline & & & \cdot \\ \hline \end{array} \right]$$

Fast Multi-Task GPs Storage and Costs for Σ^{-1} , $\log|\Sigma|$

- Assume evaluating $f(\ell, \mathbf{x})$ costs C_ℓ for any $\mathbf{x} \in [0, 1]^d$
- Assume evaluating $K((\ell, \mathbf{x}), (\ell', \mathbf{x}'))$ costs $\mathcal{O}(d)$
- Assume $m_1 \geq \dots \geq m_L$ i.e. less samples on higher levels (or reorder levels)

$$\text{NMLL} \propto \hat{\mathbf{f}}^T \Sigma^{-1} \hat{\mathbf{f}} + \log|\Sigma| + \log(2\pi)N$$

- Evaluate $\mathbf{f} = f(\mathcal{D})$ at cost $\mathcal{O}\left(\sum_{\ell=1}^L C_\ell 2^{m_\ell}\right)$
- Evaluate $\hat{\mathbf{f}} = \bar{V}\mathbf{f}$ at cost $\mathcal{O}\left(\sum_{\ell=1}^L m_\ell 2^{m_\ell}\right)$
- Evaluate only first columns of $K_{\ell\ell'}$ at total cost $\mathcal{O}\left(d \sum_{\ell=1}^L (L - \ell + 1) 2^{m_\ell}\right)$
- Evaluate Σ at cost $\mathcal{O}\left(\sum_{\ell=1}^L (L - \ell + 1) m_\ell 2^{m_\ell}\right)$
- Store Σ in $\mathcal{O}\left(\sum_{\ell=1}^L (L - \ell + 1) 2^{m_\ell}\right)$ (possibly complex) floats
- Evaluate $\Sigma^{-1} \hat{\mathbf{f}}$ and $\log|\Sigma|$ at cost $\mathcal{O}\left(\sum_{\ell'=1}^L 2^{-m_{\ell'}} \left[\sum_{\ell=1}^{\ell'-1} 2^{m_\ell}\right]^2\right)$

▶ algorithm

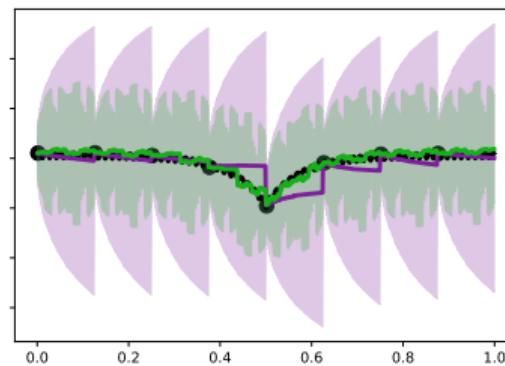
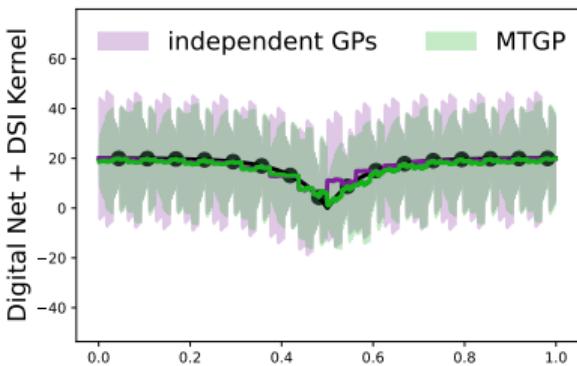
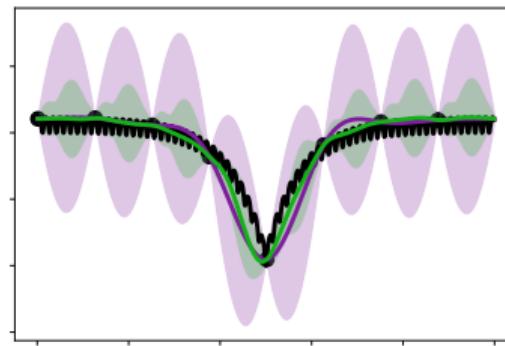
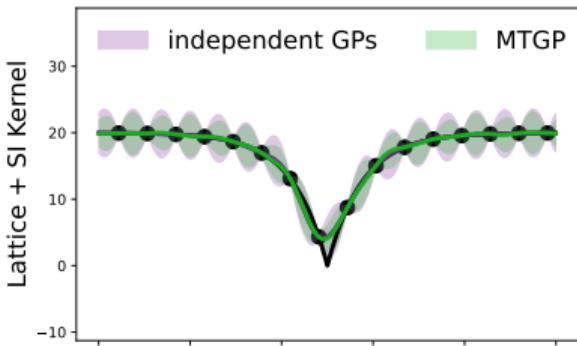
`pip install fastgps: alegresor.github.io/fastgps/`

- Single task and multi-task fast GPs (and baseline standard GPs)
- Kernel hyperparameter optimization (MLL, CV, GCV)
- Fast Bayesian cubature (including for multi-task GPs)
- Efficient batched GPR and GPU support

`pip install qmcpy: qmcsoftware.github.io/QMCSoftware/`

- Quasi-Monte Carlo Python software
- LD sequences with randomizations
 - lattices (random shifts)
 - digital nets (digital shifts, LMS, Owen scrambling (NUS), higher order nets)
 - Halton points (digital shifts, permutation scrambles, NUS)
- SI/DSI kernels of varying smoothness
- Automatic transforms (rewrite problem into $f : [0, 1]^d \rightarrow \mathbb{R}$)
- Adaptive stopping criteria (choosing N to meet user-specified error tolerance ε)
- Diverse use cases (financial options, parameterized PDEs, sensitivity indices, ...)

Thank you for listening!



References I

Michael Giles. Multilevel monte carlo path simulation. *Operations Research*, 56: 607–617, 06 2008. doi: 10.1287/opre.1070.0496.

Michael B Giles and Benjamin J Waterhouse. Multilevel quasi-monte carlo path simulation. *Advanced Financial Modelling, Radon Series on Computational and Applied Mathematics*, 8:165–181, 2009.

Jagadeeswaran Rathinavel. *Fast automatic Bayesian cubature using matching kernels and designs*. Illinois Institute of Technology, 2019.

Jagadeeswaran Rathinavel and Fred J. Hickernell. Fast automatic bayesian cubature using lattice sampling. *Statistics and Computing*, 29(6):1215–1229, Sep 2019. ISSN 1573-1375. doi: 10.1007/s11222-019-09895-9. URL
<http://dx.doi.org/10.1007/s11222-019-09895-9>.

Jagadeeswaran Rathinavel and Fred J Hickernell. Fast automatic bayesian cubature using sobol sampling. In *Advances in Modeling and Simulation: Festschrift for Pierre L'Ecuyer*, pages 301–318. Springer, 2022.

References II

S. Surjanovic and D. Bingham. Virtual library of simulation experiments: Test functions and datasets. Retrieved April 9, 2025, from
<http://www.sfu.ca/~ssurjano>.

► Return to slide on Fast Multi-Task GPs Storage and Costs

Algorithm 1 Inverse and Determinant of Σ

Require: Σ diagonal block matrix with $m_1 \geq \dots \geq m_L$.

$A \leftarrow \Sigma_{11}^{-1}$ **diagonal**

$\rho \leftarrow |\Sigma_{11}|$

$\ell \leftarrow 2$

while $\ell \leq L$ **do**

$D \leftarrow \Sigma_{\ell\ell}$ **diagonal**

$C \leftarrow \Sigma_{1:\ell-1,\ell}$ **diagonal blocks**

$S_\ell \leftarrow D - \bar{C}A^{-1}C$ **diagonal Schur complement**

$A \leftarrow \begin{pmatrix} A^{-1} + A^{-1}CS_\ell^{-1}\bar{C}A^{-1} & -A^{-1}CS_\ell^{-1} \\ -S_\ell^{-1}\bar{C}A^{-1} & S_\ell^{-1} \end{pmatrix}$

$\rho \leftarrow \rho|S_\ell|$

$\ell \leftarrow \ell + 1$

end while

return $A = \Sigma^{-1}$ and $\rho = |\Sigma|$
